

Prometheus

Prometheus

Course Introduction

Prometheus

- Thank you for taking this **Prometheus Course**
 - **Open-source** monitoring solution & time series database
 - Was built by **Soundcloud**
 - **Very active** developer and user community
 - Now its a **standalone** open source project
 - Joined the **Cloud Native Computing Foundation** in 2016
 - Ideal for monitoring on premise as well as cloud workloads

Course Overview

Introduction	Monitoring	Alerting	Internals	Use cases
What is Prometheus	Client Libraries	Introduction	Storage	Grafana Provisioning
Installing Prometheus & Grafana	Pushing metrics	Setting up alerts	Security	Cloudwatch Integration
Concepts	Querying			
Configuration	Service Discovery			
Monitoring Nodes	Exporters			
Architecture				

Course Objectives

- To be able to **use Prometheus**
- To get familiar with the **Prometheus ecosystem**
- To setup a monitoring **platform using**
 - **Prometheus**
- To **create alerts** in Prometheus
- To be able to query Prometheus data

Who is Edward Viaene

- My name is Edward Viaene
- I am a **consultant** and **trainer** in Cloud and Big Data technologies
- I'm a big advocate of **Agile** and **DevOps techniques**
- I held various roles from **banking** to **startups**
- I have a **background** in Unix/Linux, Networks, Security, Risk, and distributed computing
- Nowadays I specialize in everything that has to do with **Cloud** and **DevOps**

Who is Jorn Jambers

- My name is Jorn Jambers
- I am a **freelance DevOps consultant** and **trainer**
- DevOps **advocate**
- Worked in banks, consultancy companies and startups
 - In the latter I found my **passion** for **DevOps**
- I have a **background** in Unix/Linux, Hadoop, DBA, Networks, automations
- Today I help companies **succeed** on the public cloud

Online Training

- **Online training** on Udemy
 - **DevOps, Distributed Computing, Cloud, Big Data**
 - Using online video lectures
 - 40,000+ enrolled students in 100+ countries

Prometheus

Introduction

Prometheus

- Prometheus is an **Open source monitoring** solution
- Started at SoundCloud around 2012-2013, and was made public in early 2015
- Prometheus provides **Metrics & Alerting**
- It is inspired by Google's **Borgmon**, which uses time-series data as a datasource, to then send alerts based on this data
- It fits very well in the **cloud native infrastructure**
 - Prometheus is also a member of the **CNCF (Cloud Native Foundation)**

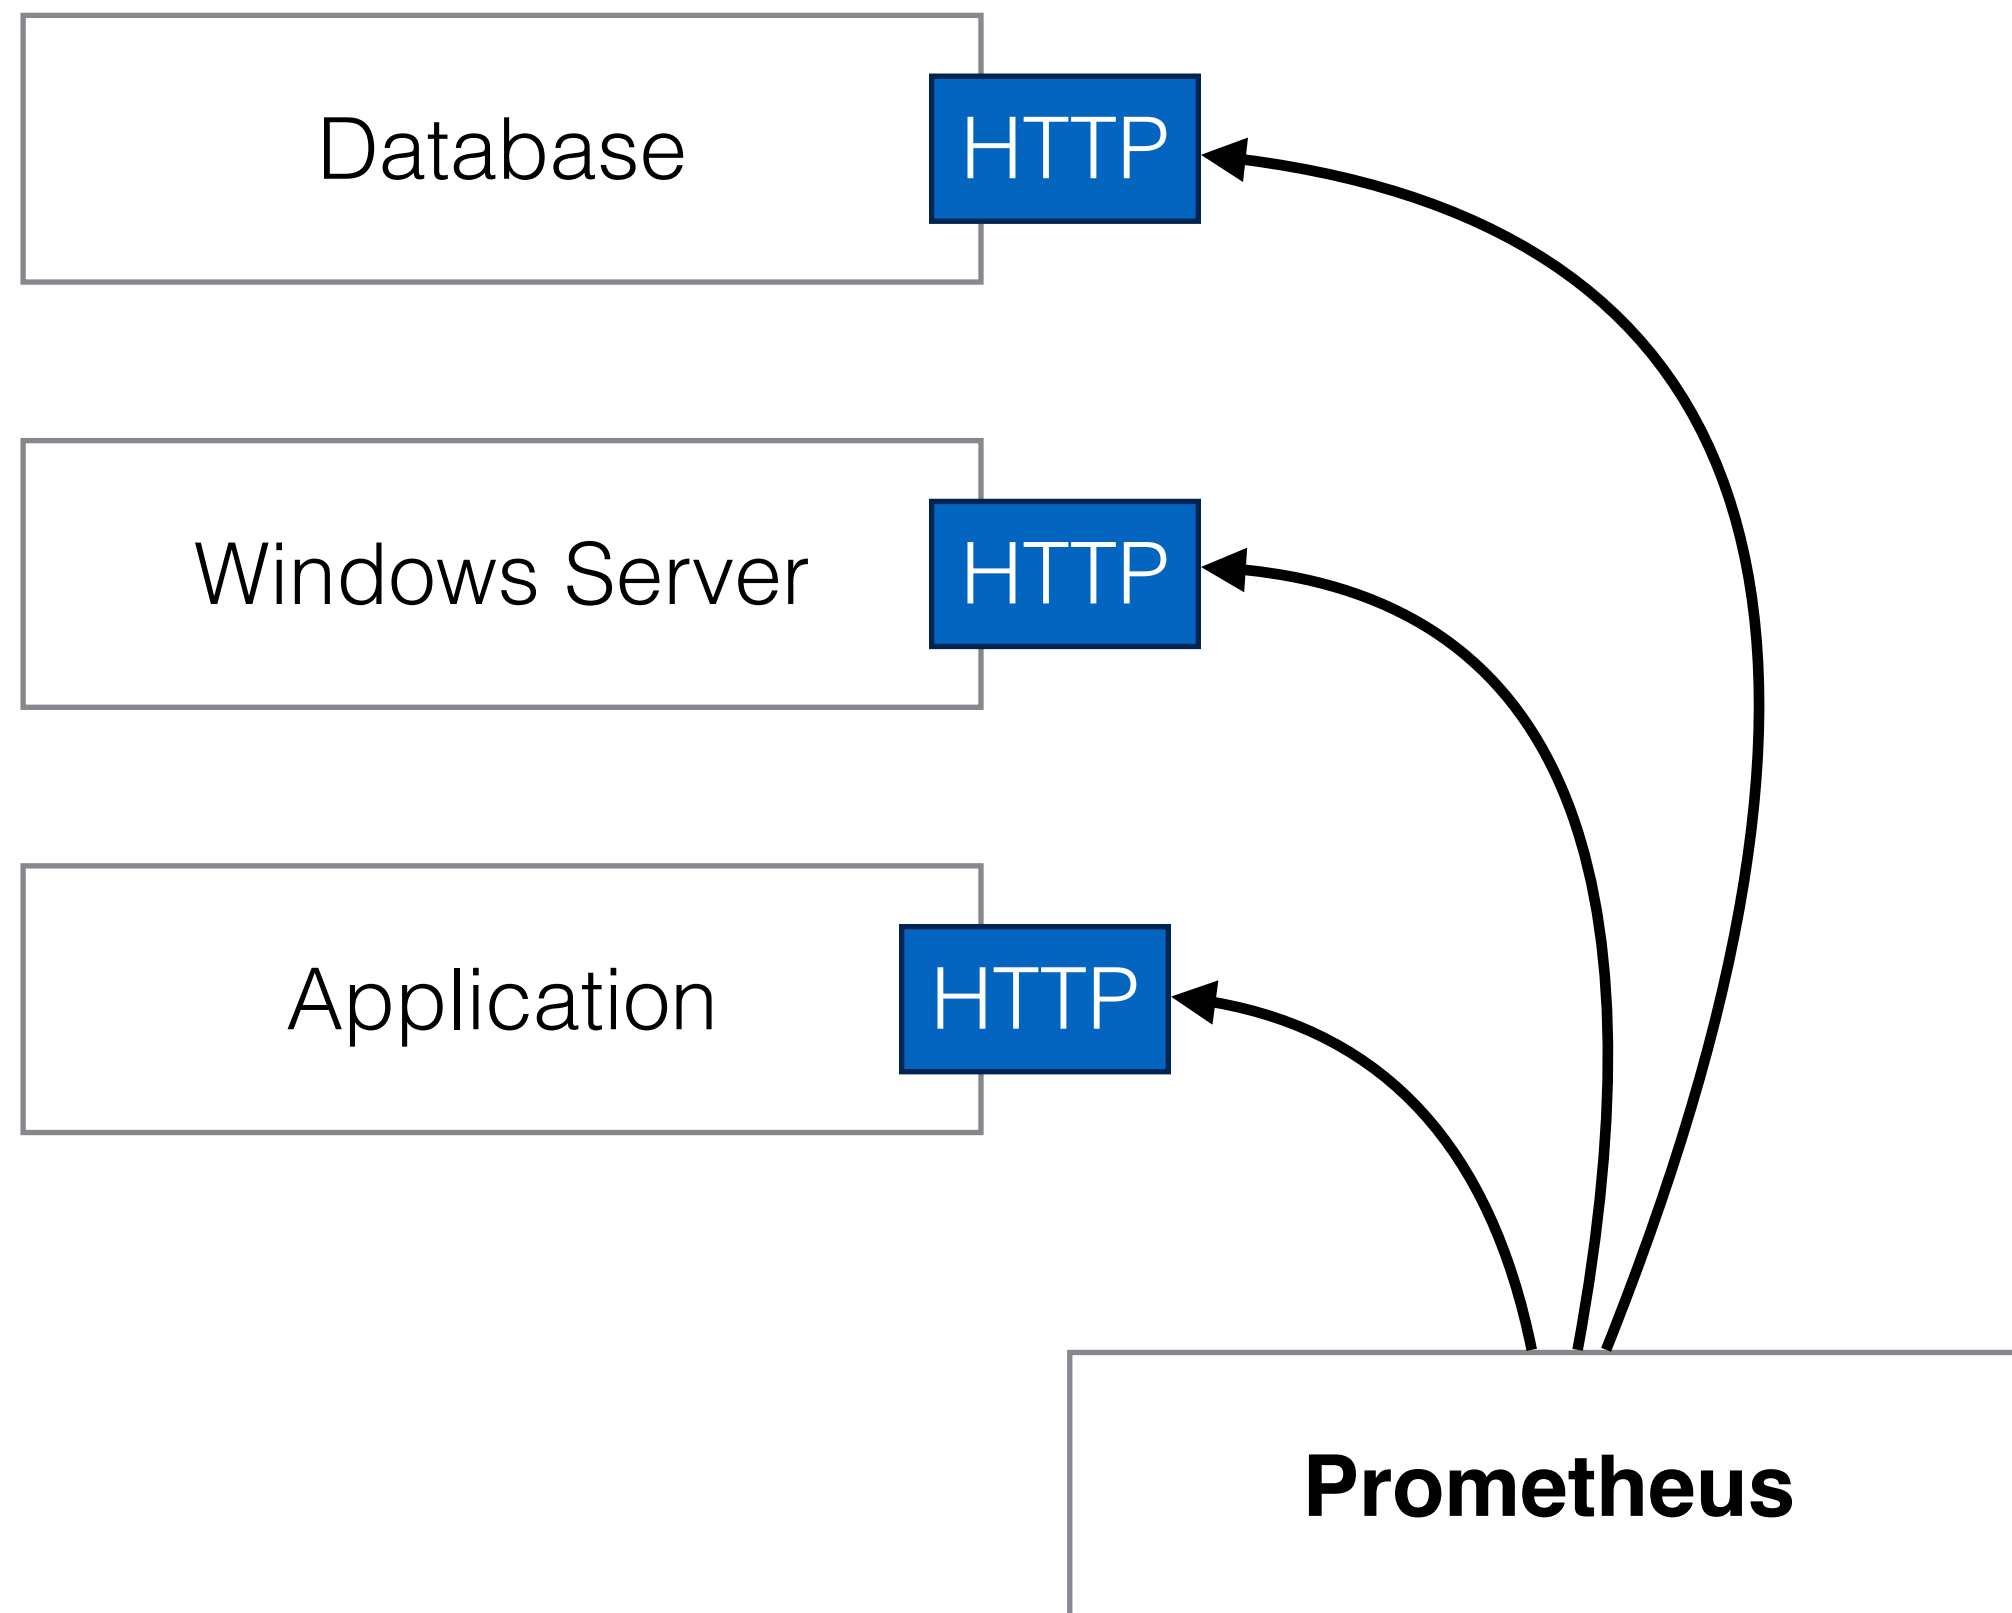
Prometheus

- In Prometheus we talk about **Dimensional Data**: time series are identified by metric name and a set of key/value pairs

Metric name	Label	Sample
Temperature	location=outside	90

- Prometheus includes a Flexible **Query Language**
- **Visualizations** can be shown using a built-in expression browser or with integrations like Grafana
- It stores metrics in **memory** and **local disk** in an own **custom, efficient format**
- It is written in **Go**
- Many **client libraries** and **integrations available**

How does Prometheus work?



- Prometheus collects metrics from monitored targets by **scraping metrics HTTP endpoints**
 - This is fundamentally different than other monitoring and alerting systems, (except this is also how Google's Borgmon works)
 - **Rather than using custom scripts** that check on particular services and systems, the **monitoring data itself is used**
- **Scraping endpoints** is much more efficient than other mechanisms, like 3rd party agents
 - A single prometheus server is able to ingest up to one million samples per second as several million time series

Prometheus

Installation

Prometheus Installation

- I will install **Prometheus using scripts** from our **GitHub** repository (<https://github.com/in4it/prometheus-course>)
 - They will work on any modern **Linux distribution**
- I'll install it on a **DigitalOcean droplet**
 - Feel free to use the scripts with any **Cloud Provider, Virtual Machine, or Docker image**, as long as it's a recent Linux distribution
 - To get a free **\$100 coupon on DigitalOcean**, valid for 60-days with a valid payment method added, use the following link:

<https://m.do.co/c/b71b388ab76f>
 - \$10 is enough to run a 2 **GB memory droplet** for one month

Prometheus Installation

- If you do not want to use the provided scripts, you can download the **full distribution** from <https://github.com/prometheus/prometheus/releases>
- **MacOS, Windows, Linux**, and some **Unix** distributions are supported
- After extracting you'll get a **prometheus executable** (prometheus.exe for windows), which you can use to run prometheus, for example:
 - `./prometheus --config.file /path/to/prometheus.yaml`
- It's best to **use the scripts we provided** so that **your environment** is the **same** as ours when you follow the demos


Demo

Installing Prometheus & Grafana

Prometheus

Basic Concepts

Concepts



- All data is stored as time series
- Every time series is identified by the “**metric name**” and a set of **key-value pairs**, called **labels**
 - **metric:** go_memstat_alloc_bytes
 - instance=localhost:9090
 - job=prometheus

Prometheus

- The time series data also consists of the **actual data**, called **Samples**:
 - It can be a **float64** value
 - or a **millisecond-precision timestamp**

Graph Console	
Element	Value
go_memstats_alloc_bytes{instance="localhost:9090",job="prometheus"}	30950392
go_memstats_alloc_bytes{instance="localhost:9100",job="node_exporter"}	2251192

Prometheus

- The notation of time series is often using this **notation**:
 - `<metric name>{<label name>=<label value>, ...}`
 - For example:
 - `node_boot_time{instance="localhost:9100",job="node_exporter"}`

Prometheus

Prometheus Configuration file

Prometheus Configuration

- The **configuration** is stored in the Prometheus configuration file, in yaml format
 - The configuration file can be **changed and applied**, without having to restart Prometheus
 - A **reload** can be done by executing `kill -SIGHUP <pid>`
- You can also pass parameters (flags) at **startup time** to `./prometheus`
 - Those parameters cannot be changed without restarting Prometheus
- The configuration file is passed using the flag `--config.file`

Prometheus Configuration

- The default configuration looks like this:

```
# my global config
global:
  scrape_interval:     15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
  evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.
  # scrape_timeout is set to the global default (10s).

# Alertmanager configuration
alerting:
  alertmanagers:
    - static_configs:
      - targets:
        # - alertmanager:9093

# Load rules once and periodically evaluate them according to the global 'evaluation_interval'.
rule_files:
  # - "first_rules.yml"
  # - "second_rules.yml"
```

Prometheus Configuration

- To scrape metrics, you need to add configuration to the prometheus config file
- For example, to scrape metrics from prometheus itself, the following code block is added by default

```
# A scrape configuration containing exactly one endpoint to scrape:
# Here it's Prometheus itself.
scrape_configs:
  # The job name is added as a label `job=<job_name>` to any timeseries scraped from this config.
  - job_name: 'prometheus'

    # metrics_path defaults to '/metrics'
    # scheme defaults to 'http'.

static_configs:
  - targets: ['localhost:9090']
```


Demo

Prometheus Configuration

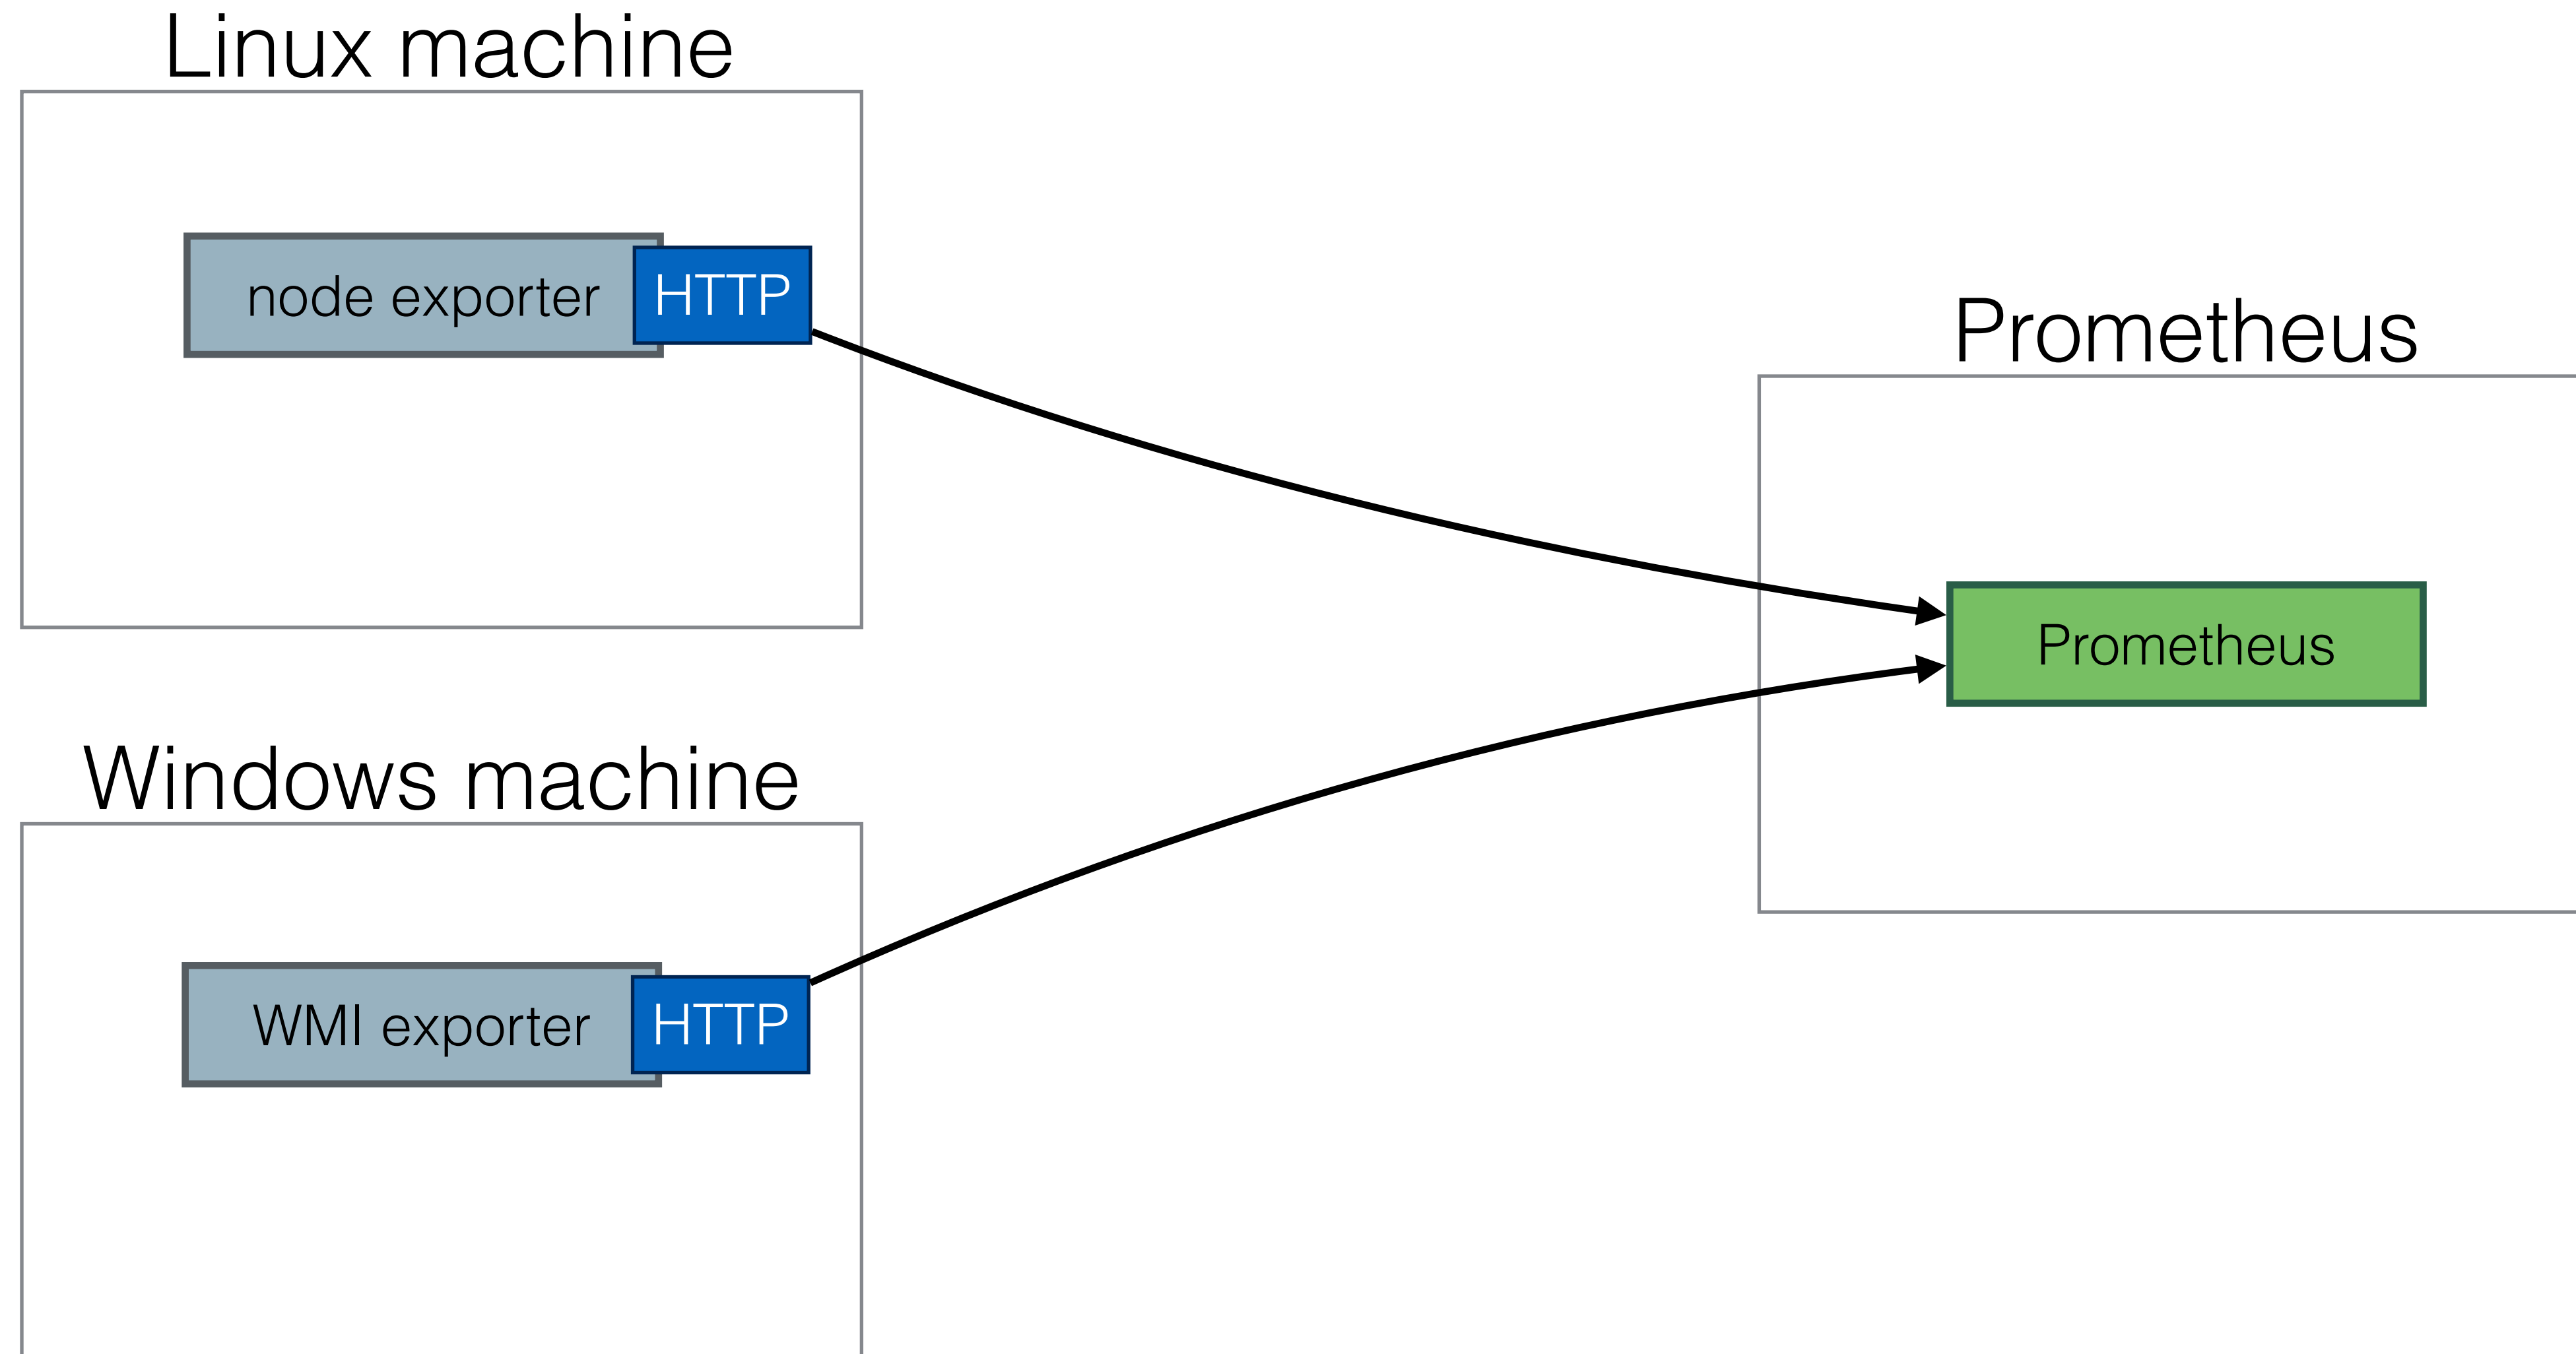
Prometheus

Monitoring Nodes

Monitor nodes

- To monitor nodes, you need to install the node-exporter
- The node exporter will expose machine metrics of Linux / *Nix machines
 - For example: cpu usage, memory usage
- The node exporter can be used to monitor machines, and later on, you can **create alerts based on these ingested metrics**
- For Windows, there's a WMI exporter (see https://github.com/martinlindhe/wmi_exporter)

Monitor nodes



Demo

Node exporter

Demo

WMI exporter

Prometheus

Monitoring

Monitoring - Introduction

- Client Libraries
- Pushing Metrics
- Querying
- Service Discovery
- Exporters

Prometheus

Client Libraries

Client Libraries - Introduction

- Instrumenting your code
- Libraries
 - Official: Go, Java/Scala, Python, Ruby
 - Unofficial: Bash, C++, Common Lisp, Elixir, Erlang, Haskell, Lua for Nginx, Lua for Tarantool, .NET / C#, Node.js, PHP, Rust
- No client library available?
 - Implement it yourself in one of the supported exposition formats

Client Libraries - Introduction

- **Exposition** formats:
 - Simple text-based format
 - Protocol-buffer format (**Prometheus 2.0 removed support for the protocol-buffer format**)

```
metric_name [
  "{" label_name "=" `"` label_value `"` { "," label_name "=" `"` label_value `"` } [ "," ] }"
] value [ timestamp ]
```

```
node_filesystem_avail_bytes{device="/dev/vda1",fstype="ext4",mountpoint="/"} 4.9386491904e+10
node_filesystem_avail_bytes{device="/dev/vda15",fstype="vfat",mountpoint="/boot/efi"} 1.05903104e+08
node_filesystem_avail_bytes{device="lxcfs",fstype="fuse.lxcfs",mountpoint="/var/lib/lxcfs"} 0
node_filesystem_avail_bytes{device="tmpfs",fstype="tmpfs",mountpoint="/run"} 2.01273344e+08
node_filesystem_avail_bytes{device="tmpfs",fstype="tmpfs",mountpoint="/run/lock"} 5.24288e+06
```

Client Libraries - Introduction

- 4 types of metrics
 - **Counter**
 - **Gauge**
 - **Histogram**
 - **Summary**

Client Libraries - Introduction

- **Counter**

A value that only goes up (e.g. Visits to a website)

Client Libraries - Introduction

- **Gauge**

Single numeric value that can go up and down (e.g. CPU load, temperature)

Client Libraries - Introduction

- **Histogram**

Samples observations (e.g. request durations or response sizes) and these observations get counted into **buckets**. Includes (`_count` and `_sum`)

Main purpose is calculating quantiles

Client Libraries - Introduction

- **Summary**

Similar to a ***histogram***, a ***summary*** samples observations (e.g. request durations or response sizes). A summary also provides a total count of observations and a sum of all observed values, it calculates configurable quantiles over a sliding time window.

Example: You need 2 counters for calculating the latency

- 1) total request(**_count**)
- 2) the total latency of those requests (**_sum**)

Take the rate() and divide = average latency

Prometheus

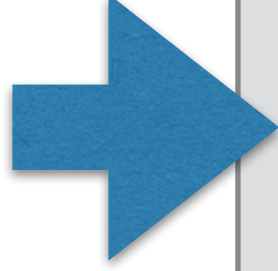
Instrumentation- Python

Client Libraries - Python Example

- https://github.com/prometheus/client_python
- Officially supported language
- `pip install prometheus_client`
- Supported metrics: Counter, Gauge, Summary and Histogram

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

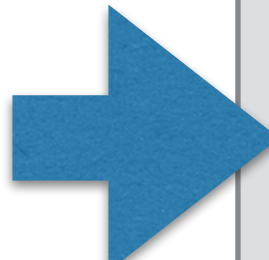
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

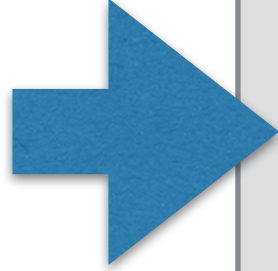
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

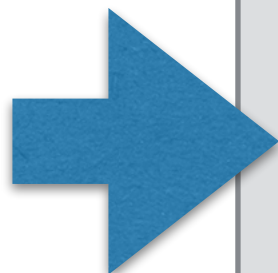
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

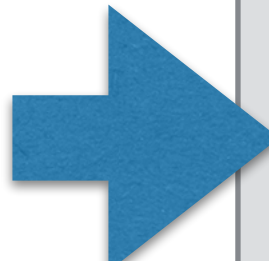
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```


Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

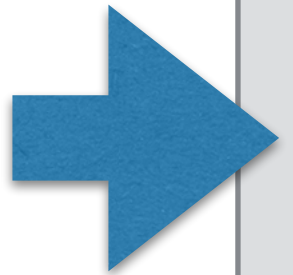
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

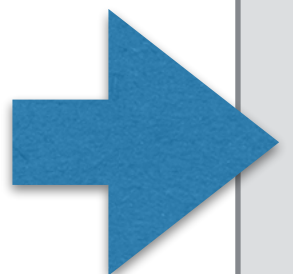
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```


Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

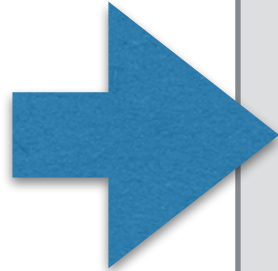
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/jorn", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:



```
import random, time

from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram

app = Flask(__name__)

REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')

@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'

@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/jorn", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)

@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

Client Libraries - Python Example

- Python example:

```
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def metrics():
    REQUESTS.labels(method='GET', endpoint="/metrics", status_code=200).inc()
    return generate_latest(REGISTRY)

if __name__ == "__main__":
    app.run(host='0.0.0.0')
```

Prometheus

Instrumentation - Go

Client Libraries - Golang Example

- https://github.com/prometheus/client_golang
- Officially supported language
- Easy to implement:

```
package main


import (
    "github.com/prometheus/client_golang/prometheus/promhttp"
    "net/http"
)

func main() {
    http.Handle("/metrics", promhttp.Handler())
    panic(http.ListenAndServe(":8080", nil))
}
```

- Supported metrics: Counter, Gauge, Summary and Histogram

Client Libraries - Golang Example

- Gauge



```
import "github.com/prometheus/client_golang/prometheus"

var jobsInQueue = prometheus.NewGauge(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs queued",
    },
)

func init(){
    prometheus.MustRegister(jobsInQueue)
}

func enqueueJob(job Job) {
    queue.Add(job)
    jobsInQueue.Inc()
}

func runNextJob() {
    job := queue.Dequeue()
    jobsInQueue.Dec()

    job.Run()
}
```

Client Libraries - Golang Example

- Gauge



```
import "github.com/prometheus/client_golang/prometheus"

var jobsQueued = prometheus.NewGauge(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs queued",
    },
)

func init(){
    prometheus.MustRegister(jobsQueued)
}

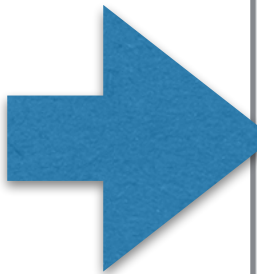
func enqueueJob(job Job) {
    queue.Add(job)
    jobsQueued.Inc()
}

func runNextJob() {
    job := queue.Dequeue()
    jobsInQueue.Dec()

    job.Run()
}
```

Client Libraries - Golang Example

- Adding labels



```
import "github.com/prometheus/client_golang/prometheus"

var jobsQueued = prometheus.NewGaugeVec(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs in the queue",
    },
    []string{"job_type"},
)

func init(){
    prometheus.MustRegister(jobsQueued)
}

func enqueueJob(job Job) {
    queue.Add(job)
    jobsInQueue.WithLabelValues(job.Type()).Inc()
}

func runNextJob() {
    job := queue.Dequeue()
    jobsInQueue.WithLabelValues(job.Type()).Dec()

    job.Run()
}
```


Client Libraries - Golang Example

- Adding labels

```
import "github.com/prometheus/client_golang/prometheus"

var jobsQueued = prometheus.NewGaugeVec(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs in the queue",
    },
    []string{"job_type"},
)

func init(){
    prometheus.MustRegister(jobsQueued)
}


func enqueueJob(job Job) {
    queue.Add(job)
    jobsQueued.WithLabelValues(job.Type()).Inc()
}

func runNextJob() {
    job := queue.Dequeue()
    jobsQueued.WithLabelValues(job.Type()).Dec()

    job.Run()
}
```

Client Libraries - Golang Example

- Histogram



```
import "github.com/prometheus/client_golang/prometheus"

var jobsDurationHistogram = prometheus.NewHistogramVec(
    prometheus.HistogramOpts{
        Name:    "jobs_duration_seconds",
        Help:    "Jobs duration distribution",
        Buckets: []float64{1, 2, 5, 10, 20, 60},
    },
    []string{"job_type"},
)

start := time.Now()
job.Run()
duration := time.Since(start)
jobsDurationHistogram.WithLabelValues(job.Type()).Observe(duration.Seconds())
```


Client Libraries - Golang Example

- Histogram

```
import "github.com/prometheus/client_golang/prometheus"

var jobsDurationHistogram = prometheus.NewHistogramVec(
    prometheus.HistogramOpts{
        Name:    "jobs_duration_seconds",
        Help:    "Jobs duration distribution",
        Buckets: []float64{1, 2, 5, 10, 20, 60},
    },
    []string{"job_type"},
)

start := time.Now()
job.Run()
duration := time.Since(start)
jobsDurationHistogram.WithLabelValues(job.Type()).Observe(duration.Seconds())
```



Client Libraries - Golang Example

- Summary



```
prometheus.NewSummary()
```


Prometheus

Pushing metrics

Pushing Metrics - Introduction

- <https://github.com/prometheus/pushgateway>

- Diagram



Pushing Metrics - Introduction

- Sometimes metrics cannot be scraped
Example: **batch jobs, servers are not reachable due to NAT, firewall**
- **Pushgateway** is used as an intermediary service which allows you to push metrics.
- **Pitfalls**
 - Most of the times this is a single instance so this results in a **SPOF**
 - Prometheus's automatic instance health monitoring is not possible
 - The **Pushgateway** never forgets the metrics unless they are deleted via the api
example:

```
curl -X DELETE http://localhost:9091/metrics/job/prom_course/instance/localhost
```

Pushing Metrics - Introduction

- Only 1 valid use case for the Pushgateway
 - Service-level batch jobs and not related to a specific machine
- If NAT or/both firewall is blocking you from using the pull mechanism
 - Move the Prometheus server on the same network

Pushing Metrics - Python Example

- Python example:

```
from prometheus_client import CollectorRegistry, Gauge, push_to_gateway

registry = CollectorRegistry()
g = Gauge('job_last_success_unixtime', 'Last time the course batch job has finished', registry=registry)
g.set_to_current_time()
push_to_gateway('localhost:9091', job='batchA', registry=registry)
```

- **Pushgateway** functions take a grouping key.
- **push_to_gateway** replaces metrics with the same grouping key
- **pushadd_to_gateway** only replaces metrics with the same name and grouping key
- **delete_from_gateway** deletes metrics with the given job and grouping key.

Prometheus

Pushing Metrics - Go

Pushing Metrics - Go Example

- Go example:

```
package main
import (
    "flag"
    "log"
    "net/http"
    "github.com/prometheus/client_golang/prometheus/promhttp"
    "github.com/prometheus/client_golang/prometheus/push"
)

gatewayUrl:="http://localhost:9091/"

throughputGauge := prometheus.NewGauge(prometheus.GaugeOpts{
    Name: "throughput",
    Help: "Throughput in Mbps",
})
throughputGauge.Set(800)

if err := push.Collectors(
    "throughput_job", push.HostnameGroupingKey(),
    gatewayUrl, throughputGauge
); err != nil {
    fmt.Println("Could not push completion time to Pushgateway:", err)
}
```

Prometheus

Querying

Querying Metrics - Introduction

- Prometheus provides a functional expression language called PromQL
 - Provides built in operators and functions
 - Vector-based calculations like Excel
 - Expressions over time-series vectors
- PromQL is read-only
- Example:

```
100 - (avg by (instance) (irate(node_cpu_seconds_total{job='node_exporter',mode="idle"}[5m])) * 100)
```

Prometheus

Querying - Expressions

Querying Metrics - Introduction

- **Instant vector** - a set of time series containing a single sample for each time series, all sharing the same timestamp
Example: `node_cpu_seconds_total`
- **Range vector** - a set of time series containing a range of data points over time for each time series
Example: `node_cpu_seconds_total[5m]`
- **Scalar** - a simple numeric floating point value
Example: `-3.14`
- **String** - a simple string value; currently unused
Example: `foobar`

Prometheus

Querying - Operators

Querying Metrics - Introduction

- **Arithmetic binary operators**

Example: - (subtraction), * (multiplication), / (division), % (modulo), ^ (power/exponentiation)

- **Comparison binary operators**

Example: == (equal), != (not-equal), > (greater-than), < (less-than), >= (greater-or-equal), <= (less-or-equal)

- **Logical/set binary operators**

Example: and (intersection), or (union), unless (complement)

- **Aggregation operators**

Example: **sum** (calculate sum over dimensions), **min** (select minimum over dimensions), **max** (select maximum over dimensions), **avg** (calculate the average over dimensions), **stddev** (calculate population standard deviation over dimensions), **stdvar** (calculate population standard variance over dimensions), **count** (count number of elements in the vector), **count_values** (count number of elements with the same value), **bottomk** (smallest k elements by sample value), **topk** (largest k elements by sample value), **quantile** (calculate ϕ -quantile ($0 \leq \phi \leq 1$) over dimensions)

Demo

Querying

Prometheus

Service Discovery

Service Discovery - Introduction

- Definition:
Service discovery is the automatic detection of devices and services offered by these devices on a computer network.

- Not really a service discovery mechanism

```
static_configs:  
  - targets: ['localhost:9090']
```

- Cloud support for (AWS, Azure, Google,...)
- Cluster managers (Kubernetes, Marathon, ...)
- Generic mechanisms (DNS, Consul, Zookeeper, ...)

Prometheus

Service Discovery - Example AWS

Service Discovery - Introduction

- EC2 Example:

Add following config to */etc/prometheus/prometheus.yml*

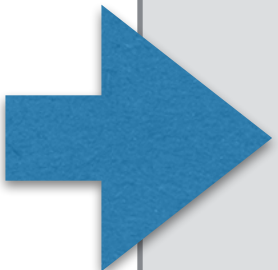
```
global:
  scrape_interval: 1s
  evaluation_interval: 1s

scrape_configs:
  - job_name: 'node'
    ec2_sd_configs:
      - region: eu-west-1
        access_key: PUT_THE_ACCESS_KEY_HERE
        secret_key: PUT_THE_SECRET_KEY_HERE
        port: 9100
```

- Make sure the user has the following IAM role: **AmazonEC2ReadOnlyAccess**
- Make sure you security groups allow access to port (9100, 9090)

Service Discovery - Introduction

- EC2 Example:
Only monitor instances started with the name PROD



```
global:
  scrape_interval: 1s
  evaluation_interval: 1s


scrape_configs:
  - job_name: 'node'
    ec2_sd_configs:
      - region: eu-west-1
        access_key: PUT_THE_ACCESS_KEY_HERE
        secret_key: PUT_THE_SECRET_KEY_HERE
        port: 9100
    relabel_configs:
      # Only monitor instances with a tag Name starting with "PROD"
      - source_labels: [__meta_ec2_tag_Name]
        regex: PROD.*
        action: keep
      # Use the instance ID as the instance label
      - source_labels: [__meta_ec2_instance_id]
        target_label: instance
```

Service Discovery - Introduction

- EC2 Example:
Relabel ip address to instance id for convenience

```
global:
  scrape_interval: 1s
  evaluation_interval: 1s

scrape_configs:
  - job_name: 'node'
    ec2_sd_configs:
      - region: eu-west-1
        access_key: PUT_THE_ACCESS_KEY_HERE
        secret_key: PUT_THE_SECRET_KEY_HERE
        port: 9100
    relabel_configs:
      # Only monitor instances with a tag Name starting with "PROD"
      - source_labels: [__meta_ec2_tag_Name]
        regex: PROD.*
        action: keep
      # Use the instance ID as the instance label
      - source_labels: [__meta_ec2_instance_id]
        target_label: instance
```



Prometheus

Service Discovery - Example Kubernetes

Service Discovery - Introduction

- Kubernetes Example:
Add following config to */etc/prometheus/prometheus.yml*

```
- job_name: 'kubernetes'
  kubernetes_sd_configs:
  -
    api_servers:
      - https://kubernetes.default.svc

    in_cluster: true

    basic_auth:
      username: prometheus
      password: secret
      retry_interval: 5s

- job_name: 'kubernetes-service-endpoints'
  kubernetes_sd_configs:
  -
    api_servers:
      - https://kube-master.prometheuscourse.com

    in_cluster: true
```

Prometheus

Service Discovery - Example DNS

Service Discovery - DNS

- DNS Example:

Add following config to */etc/prometheus/prometheus.yml*

```
- job_name: mysql
  dns_sd_configs:
    - names:
      - metrics.mysql.example.com
- job_name: haproxy
  dns_sd_configs:
    - names:
      - metrics.haproxy.example.com
```

Prometheus

Service Discovery - Example using file

Service Discovery - Using file

- File Example:

Add following config to */etc/prometheus/prometheus.yml*

```
scrape_configs:
  - job_name: 'dummy' # This is a default value, it is
    mandatory.
    file_sd_configs:
      - files:
        - targets.json
```

- Format target.json

```
[
  {
    "targets": [ "myslave1:9104", "myslave2:9104" ],
    "labels": {
      "env": "prod",
      "job": "mysql_slave"
    }
  },
  {
    "targets": [ "mymaster:9104" ],
    "labels": {
      "env": "prod",
      "job": "mysql_master"
    }
  }
]
```

Prometheus

Exporters

Exporters - Introduction

- Build for exporting prometheus metrics from existing third-party metrics
- When Prometheus is not able to pull metrics directly(Linux sys stats, haproxy, ...)
- Examples:
 - MySQL server exporter
 - Memcached exporter
 - Consul exporter
 - Node/system metrics exporter
 - MongoDB
 - Redis
 - Many more....
- <https://prometheus.io/docs/instrumenting/exporters/>

Exporters - Introduction

- We are already using one :-)
Check */etc/prometheus/prometheus.yml*

```
- job_name: 'node_exporter'  
  scrape_interval: 5s  
  static_configs:  
    - targets: ['localhost:9100']
```

Prometheus

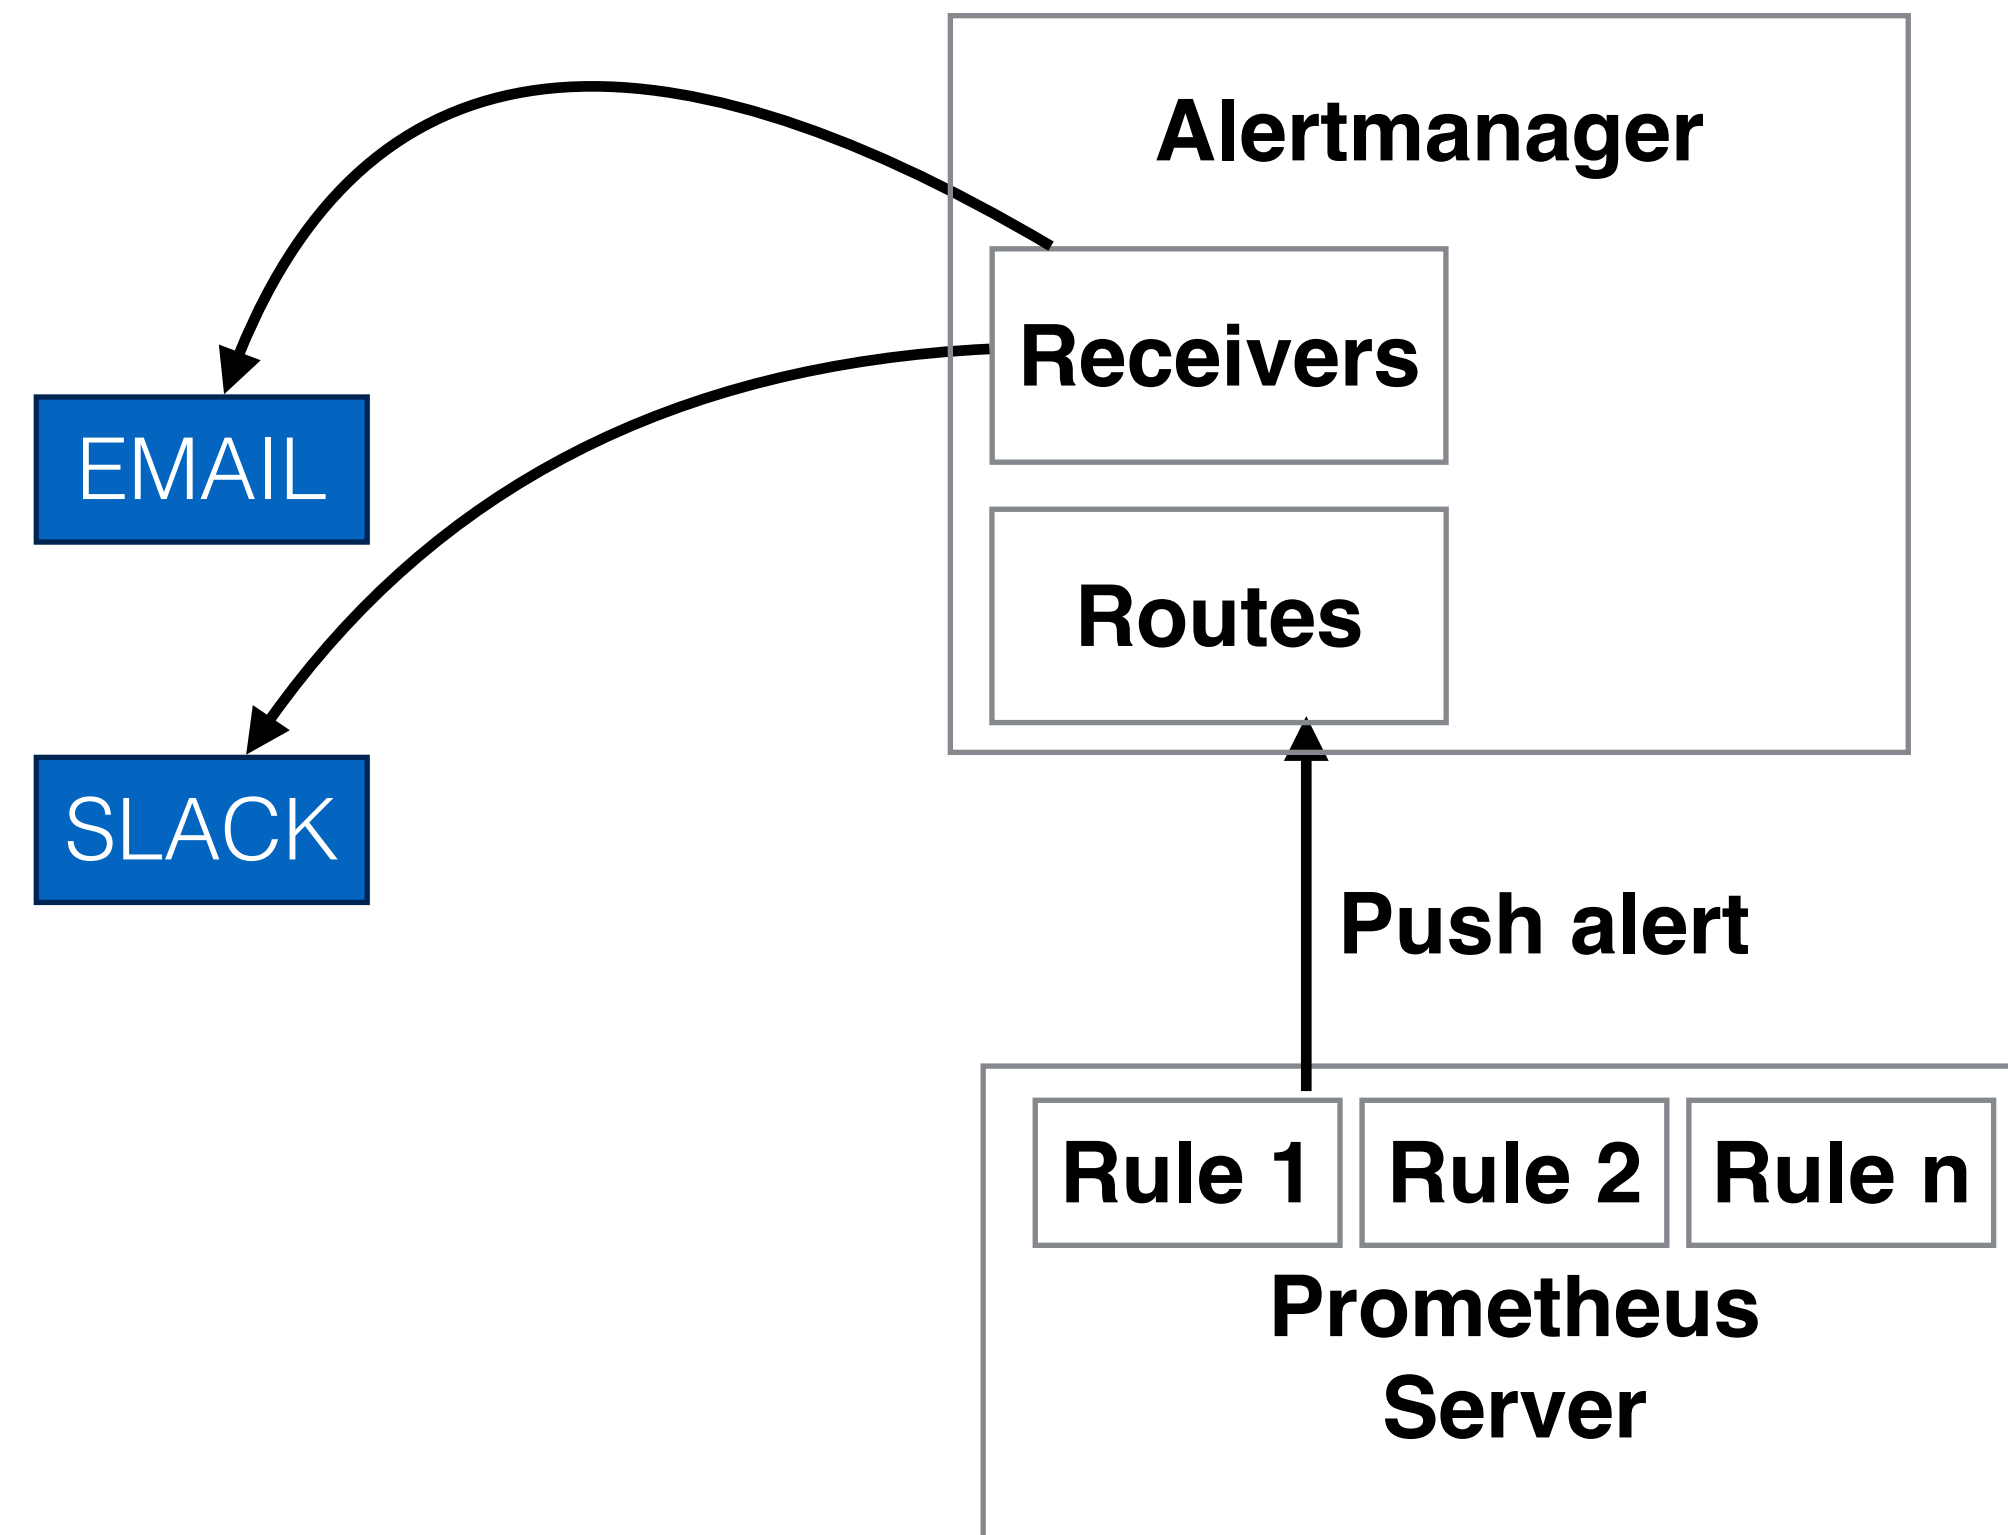
Alerting

Prometheus

Alerting - Introduction

Alerting - Introduction

- Alerting in Prometheus is separated into 2 parts
 - **Alerting rules** in Prometheus server
 - **Alertmanager**



Prometheus

Alerting - Alerting rules

Alerting Rules

- Rules live in Prometheus server config
- Best practice to separate the alerts from the prometheus config

- Add an include to */etc/prometheus/prometheus.yml*

```
rule_files:  
- "/etc/prometheus/alert.rules"
```

- Alert format:

```
ALERT <alert name>  
IF <expression>  
[ FOR <duration> ]  
[ LABELS <label set> ]  
[ ANNOTATIONS <label set> ]
```

- Alert example:

```
groups:  
- name: example  
  rules:  
  - alert: cpuUsage  
    expr: 100 - (avg by (instance) (irate(node_cpu_seconds_total{job='node_exporter',mode='idle'}[5m])) * 100) >  
    95  
    for: 1m  
    labels:  
      severity: critical  
    annotations:  
      summary: Machine under heavy load
```

Alerting Rules

- Alerting rules allow you to define the alert conditions
- Alerting rules sent the alerts being fired to an external service
- The format of these alerts is in the Prometheus expression language

- Example:

```
groups:  
- name: Important instance  
  rules:  
  
# Alert for any instance that is unreachable for >5 minutes.  
- alert: InstanceDown  
  expr: up == 0  
  for: 5m  
  labels:  
    severity: page  
  annotations:  
    summary: "Instance {{ $labels.instance }} down"  
    description: "{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 5 minutes."
```

Prometheus

Alerting - Alertmanager

Alertmanager

- **Alertmanager** handles the alerts fired by the prometheus server
- Handles **deduplication**, **grouping** and **routing** of alerts
- **Routes** alerts to **receivers** (Pagerduty, Opsgenie, email, Slack,...)

Alertmanager

- Alertmanager Configuration (*/etc/alertmanager/alertmanager.yml*):

```
global:
  smtp_smarthost: 'localhost:25'
  smtp_from: 'alertmanager@prometheus.com'
  smtp_auth_username: ''
  smtp_auth_password: ''

templates:
- '/etc/alertmanager/template/*.tmpl'

route:
  repeat_interval: 1h
  receiver: operations-team

receivers:
- name: 'operations-team'
  email_configs:
  - to: 'operations-team+alerts@example.org'
  slack_configs:
  - api_url: https://hooks.slack.com/services/XXXXXX/XXXXXX/XXXXXX
    channel: '#prometheus-course'
    send_resolved: true
```

Alertmanager

- Prometheus Configuration (*/etc/prometheus/prometheus.yml*):

```
# my global config
global:
  scrape_interval:      15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
  evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.
  # scrape_timeout is set to the global default (10s).

# Alertmanager configuration
alerting:
  alertmanagers:
    - static_configs:
      - targets:
        - localhost:9093

# Load rules once and periodically evaluate them according to the global 'evaluation_interval'.
rule_files:
  # - "first_rules.yml"
  # - "second_rules.yml"

# A scrape configuration containing exactly one endpoint to scrape:
# Here it's Prometheus itself.
scrape_configs:
  # The job name is added as a label `job=<job_name>` to any timeseries scraped from this config.
  - job_name: 'prometheus'

    # metrics_path defaults to '/metrics'
    # scheme defaults to 'http'.

    static_configs:
      - targets: ['localhost:9090']
```

Alertmanager

- Concepts:
 - **Grouping**: Groups similar alerts into 1 notification
 - **Inhibition**: Silence other alerts if one specified alert is already fired
 - **Silences**: A simple way to mute certain notifications

Alertmanager

- **High availability**
 - You can create a high available **Alertmanager** cluster using mesh config
 - Do **not** load balance this service!
 - Use a list of Alertmanager nodes in Prometheus config
 - All alerts are sent to all known Alertmanager nodes
 - No need to monitor the monitoring
 - Guarantees the notification is at least send once

Alertmanager

- **Alert states:**

Inactive - No rule is met

Pending - Rule is met but can be suppressed due to validations

Firing - Alert is sent to the configured channel(mail,Slack,...)

- Runs on port **:9093**

The screenshot shows the Alertmanager web interface. At the top, there is a navigation bar with links for 'Alertmanager', 'Alerts', 'Silences', and 'Status'. A 'New Silence' button is located on the right. Below the navigation bar, there is a section for filtering alerts. It includes tabs for 'Filter' and 'Group'. To the right of these tabs, there are radio buttons for 'Receiver: All', 'Silenced', and 'Inhibited'. Below the tabs, there is a text input field for a custom matcher, with a placeholder text 'Custom matcher, e.g. env="production"'. A blue button with a '+' sign is next to the input field. At the bottom of the interface, there is a yellow banner that says 'No alerts found'.

Alertmanager

- **Notifying multiple destinations**

```
route:
  repeat_interval: 1h
  receiver: operations-team

receivers:
- name: 'operations-team'
  email_configs:
  - to: 'operations-team+alerts@example.org'
  slack_configs:
  - api_url: https://hooks.slack.com/services/XXXXXX/XXXXXX/XXXXXX
    channel: '#prometheus-course'
    send_resolved: true
```

Prometheus

Setting up alerts

Prometheus

Setting up alerts - Demo

Setting up alerts

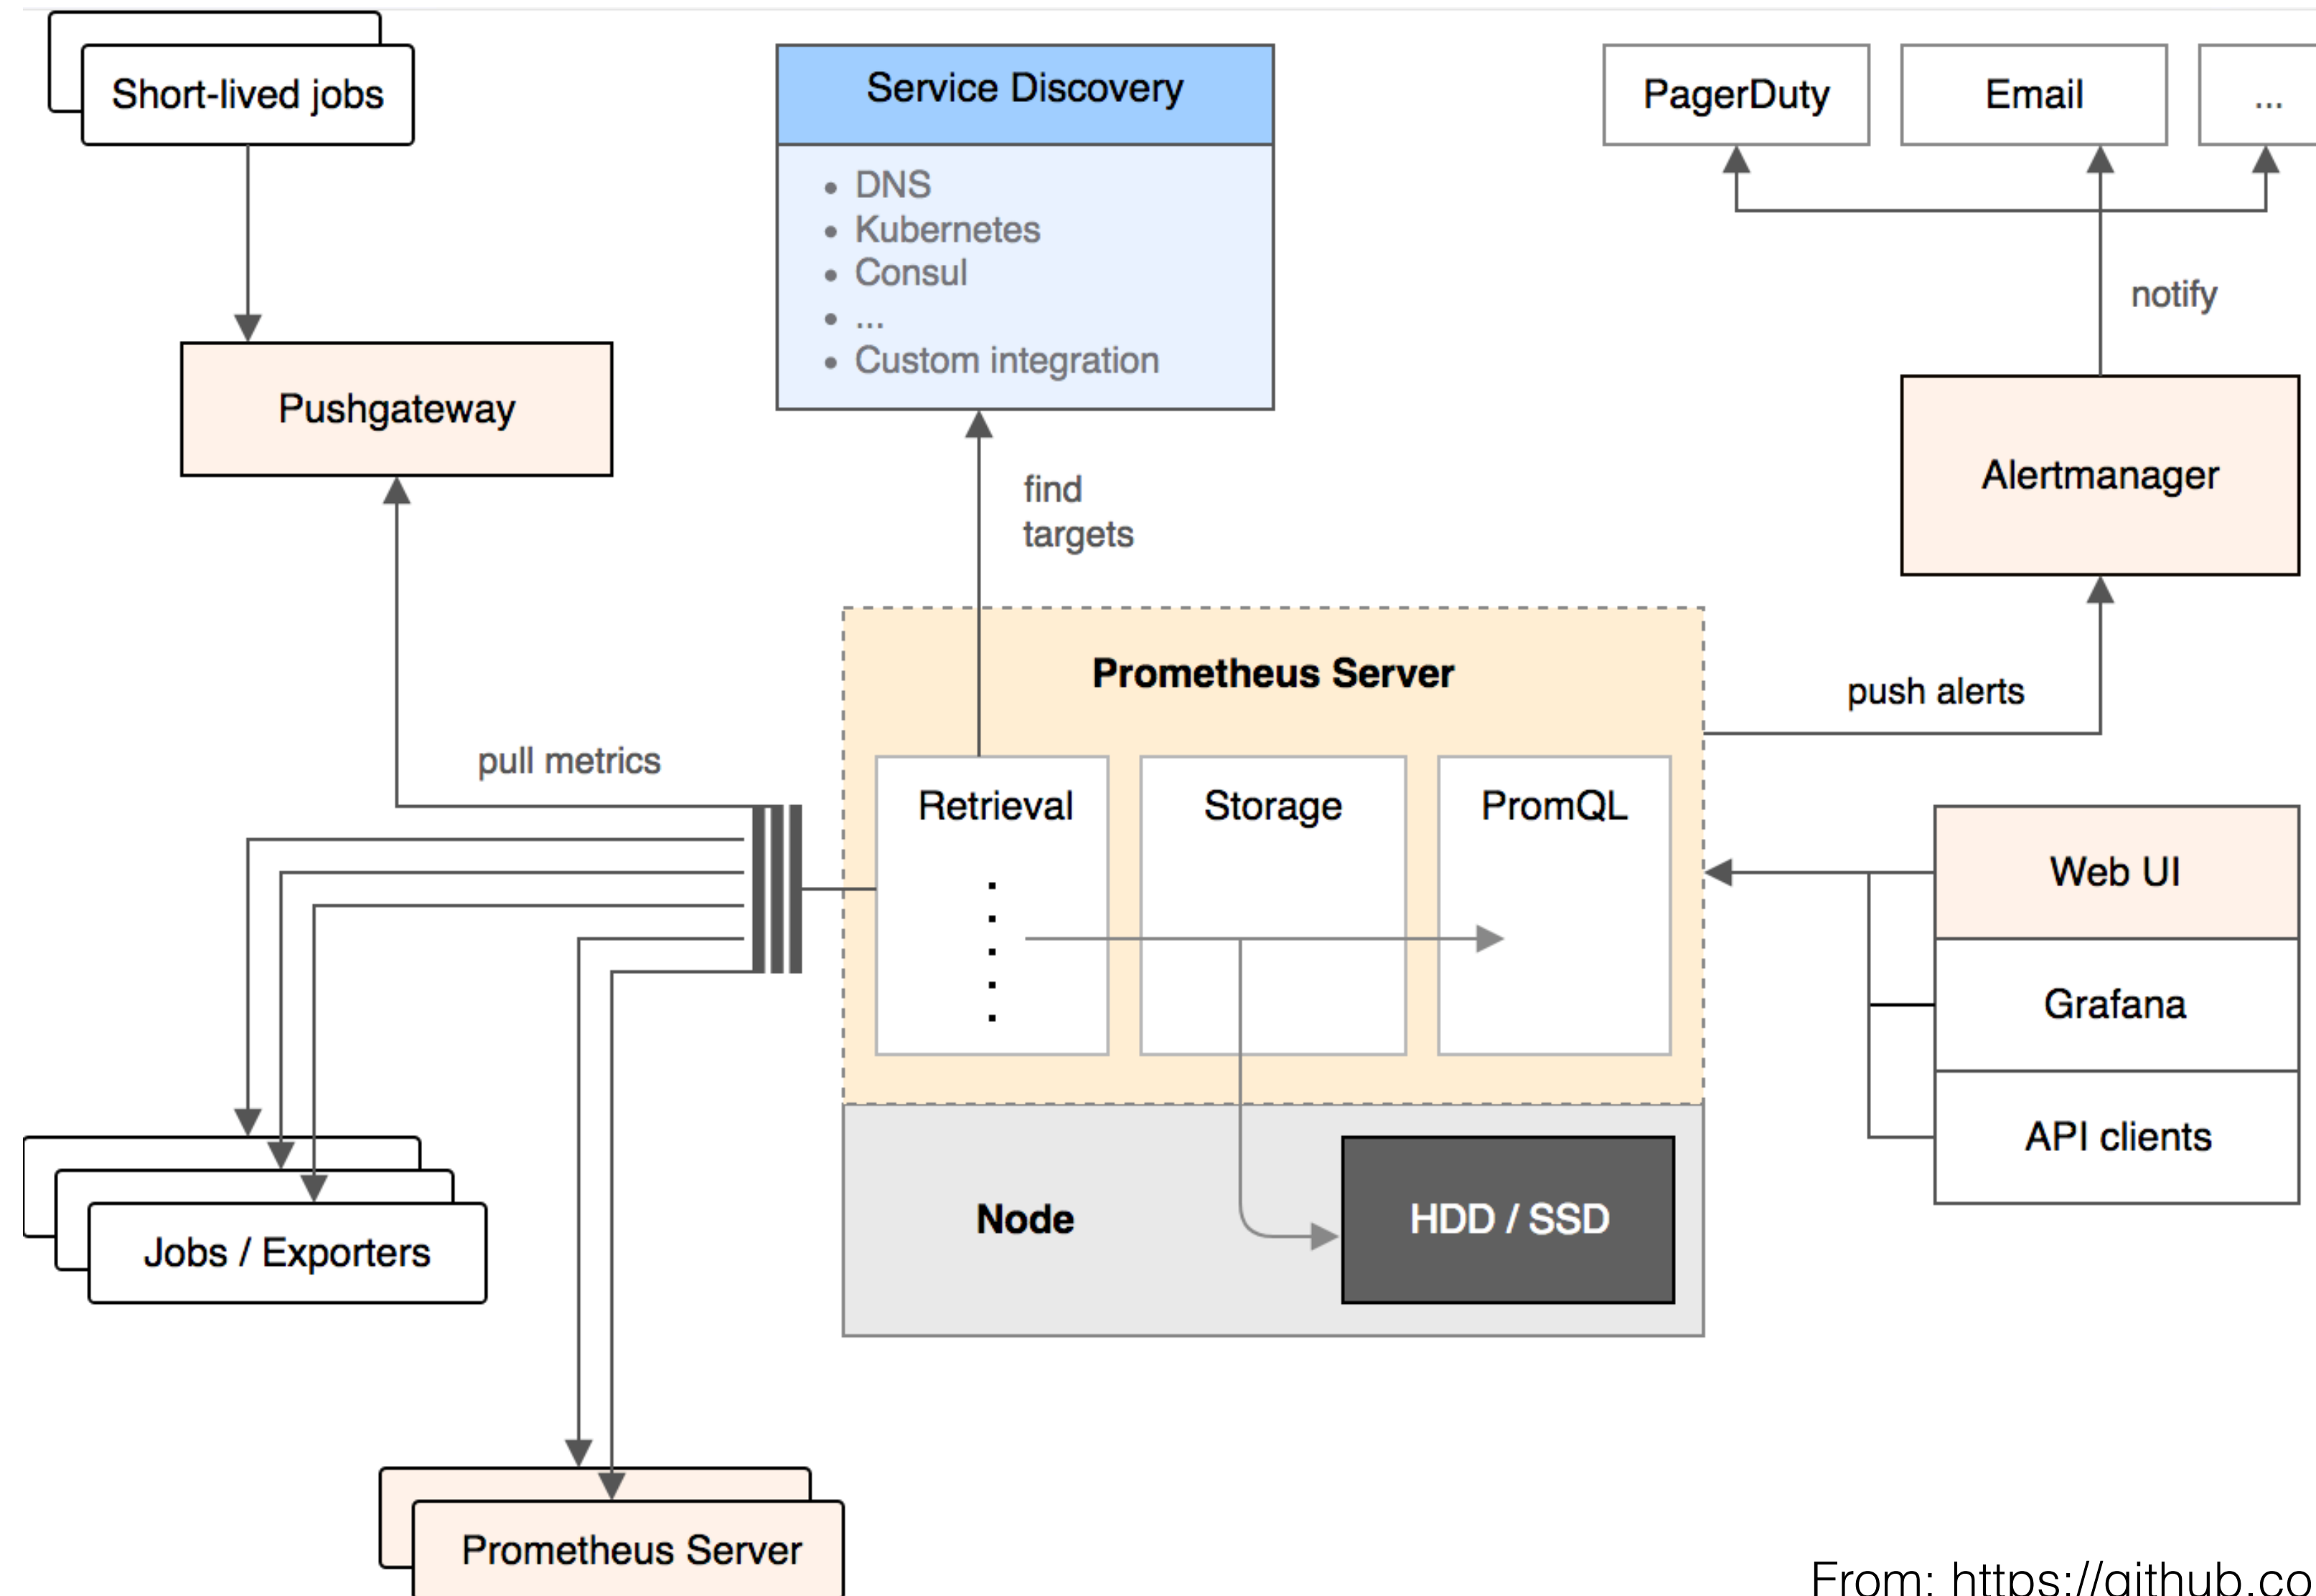
- Install Alertmanager
- Create config for the Alertmanager
 - Mail
 - Slack
- Alter prometheus config
- Setup an alert
 - See the notification coming in when an alert is fired

Prometheus internals

Prometheus

Architecture

Architecture



From: <https://github.com/prometheus/prometheus>

Prometheus

Storage

Storage

- You can use the default local **on-disk storage**, or optionally the **remote storage system**
 - **Local storage**: a local **time series database** in a custom Prometheus format
 - **Remote storage**: you can read/write samples to a **remote system** in a **standardized format**
 - Currently it uses a **snappy-compressed protocol buffer encoding** over HTTP, but might change in the future (to use gRPC or HTTP/2)

Remote Storage

- Remote storage is primarily focussed at long term storage
- Currently there are adapters available for the following solutions:

AppOptics: write	Graphite: write
Chronix: write	InfluxDB: read and write
Cortex: read and write	OpenTSDB: write
CrateDB: read and write	PostgreSQL/TimescaleDB: read and write
Gnocchi: write	SignalFx: write

Source: <https://prometheus.io/docs/operating/integrations/#remote-endpoints-and-storage>

Local Storage

- Prometheus ≥ 2.0 uses a **new storage engine** which dramatically increases scalability
- Ingested samples are **grouped in blocks of two hours**
- Those 2h samples are stored in **separate directories** (in the data directory of prometheus)
- Writes are batched and written to disk in **chunks**, containing multiple data points

directory 1

2h of data:

chunks/000001
chunks/000002

directory 2

2h of data:

chunks/000001

directory 3

2h of data:

chunks/000001
chunks/000002

Local Storage

- Every directory also has an **index file** (index) and a **metadata file** (meta.json)
- It stores the **metric names** and the **labels**, and **provides an index** from the metric names and labels to the series in the chunk files

directory 1

chunks/000001
chunks/000002
meta.json
index

directory 2

chunks/000001
meta.json
index

directory 3

chunks/000001
chunks/000002
meta.json
index

Local Storage

- The most recent data is kept **in memory**
- You don't want to loose the in-memory data during a crash, so the **data** also **needs to be persisted to disk**. This is done using a **write-ahead-log** (WAL)

directory 1

```
chunks/000001
chunks/000002
meta.json
index
```

directory 2

```
chunks/000001
meta.json
index
```

directory 3

```
chunks/000001
chunks/000002
meta.json
index
```

wal:

```
000001
000002
```

Local Storage

- **Write Ahead Log** (WAL)
 - It's quicker to **append** to a file (like a log) than making (multiple) random read/writes
 - If there's a server crash and the data from memory is lost, then the WAL will be **replayed**
 - This way, **no data will be lost** or corrupted during a crash

Local Storage

- When series gets deleted, a **tombstone file** gets created
- This is **more efficient** than immediately deleting the data from the chunk files, as the **actual delete can happen at a later time** (e.g. when there's not a lot of load)

directory 1

2h of data:
chunks/000001
chunks/000002
meta.json
index
tombstone

directory 2

2h of data:
chunks/000001
chunks/000002
meta.json
index
tombstone

directory 3

2h of data:
chunks/000001
chunks/000002
meta.json
index
tombstone

Local Storage

- The initial 2-hour blocks are **merged** in the **background** to form longer blocks
- This is called **compaction**

directory 1

2h of data:

chunks/000001
chunks/000002

directory 2+3

4h of data:

chunks/000001
chunks/000002

Local Storage

- **Block characteristics:**
 - A block on the filesystem is a **directory with chunks**
 - You can see each block as a **fully independent database** containing all time series for the window
 - Every block of data, except the current block, is **immutable** (no changes can be made)
 - These non-overlapping blocks are actually a **horizontal partitioning** of the ingested time series data

Local Storage

- This **horizontal partitioning** gives a lot of benefits:
 - When querying, the **blocks not in the time range** can be **skipped**
 - When **completing a block**, data only needs to be **added**, and not modified (avoids write-amplification)
 - **Recent data is kept in memory**, so can be queried quicker
 - **Deleting old data** is only a matter of **deleting directories** on the filesystem

Local Storage

- **Compaction:**
 - **When querying**, blocks have to be **merged** together to be able to **calculate the results**
 - Too many blocks could cause **too much merging overhead**, so blocks are compacted
 - 2 blocks are merged and **form a newly created** (often larger) **block**
 - Compaction can also **modify data: dropping deleted data** or **restructuring the chunks** to increase the query performance

Local Storage

- The index:
 - Having **horizontal partitioning** already makes **most queries quicker**, but not those that need to go through all the data to get the result
 - The index is an **inverted index** to provide better query performance, also in cases where all data needs to be queried
 - Each series is assigned a unique ID (e.g. ID 1, 2, and 3)
 - The index will contain an inverted index for the labels, for example for label env=production, it'll have 1 and 3 as IDs if those series contain the label env=production

Local Storage

- What about Disk size?
- On average, Prometheus needs **1-2 bytes per sample**
- You can use the following formula to calculate the disk space needed:

```
needed_disk_space = retention_time_seconds * ingested_samples_per_second * bytes_per_sample
```

Local Storage

- How to reduce disk size?
 - You can increase the scrape interval, which will get you less data
 - You can decrease the targets or series you scrape
 - Or you can reduce the retention (how long you keep the data)

`--storage.tsdb.retention`: This determines when to remove old data. Defaults to 15d.

References

- To read the full story of Prometheus time series database, read the blog post from **Fabian Reinartz** at <https://fabxc.org/tsdb/>

Prometheus

Security

Security

- At the moment Prometheus **doesn't offer any support for authentication** or encryption (TLS) on the server components
 - They argue that they're **focussing on building a monitoring solution**, and want to avoid having to implement complex security features
 - You can still enable **authentication** and **TLS**, using a reverse proxy
- This is **only valid for server components**, prometheus can **scrape TLS and authentication** enabled targets
 - See **tls_config** in the prometheus configuration to configure a CA certificate, user certificate and user key
 - You'd still need to setup a reverse proxy for the targets itself

Demo

Prometheus TLS and authentication

Demo

Prometheus mutual TLS for targets

Prometheus Use Cases

Monitoring a web app

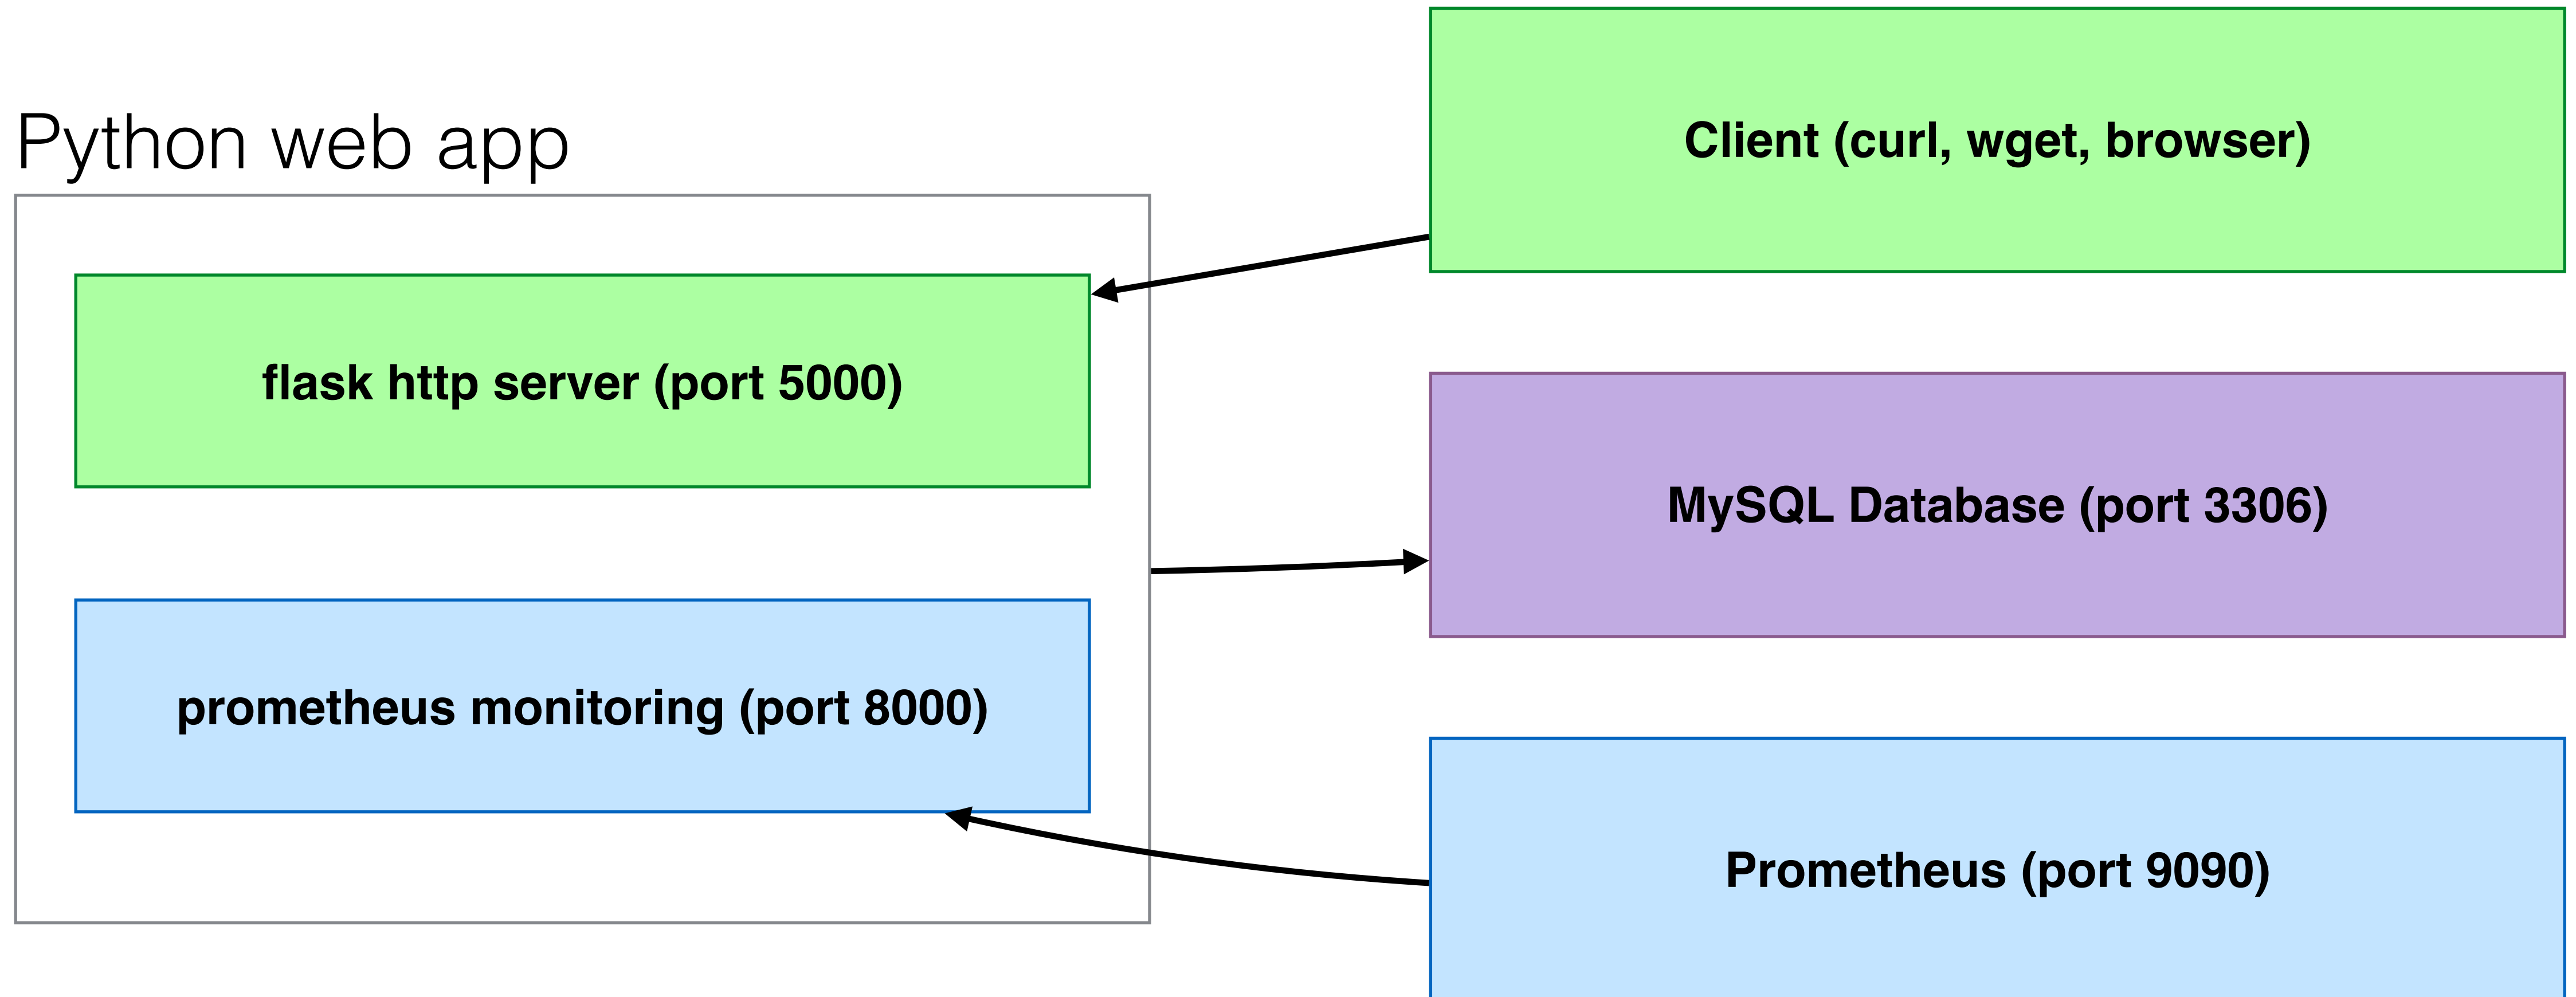
Prometheus with python-flask and MySQL

Monitoring a web app

- I'm going to integrate **prometheus monitoring** with a **web application** based on **python**
 - I'll use the official **prometheus_client** library for Python
 - **Flask** is the web framework I'm going to use
 - It will create an **http server** and I'll be able to configure routes (e.g. /query)
 - I'll use **mysqlclient** for python to query a **MySQL** database
 - I'll include one normal query and one “**bad behaving**” **query** that will take between 0 and 10 seconds to execute

The web app

Python web app



Monitoring a web app

- I'm going to use the **Counter** and the **Histogram metric types** to capture the data:
 - A **Counter** to capture the amount of times an **http endpoint** is hit + to capture the amount of times a **MySQL query** is executed
 - The value of the **Counter must always increase**, that's why you should take the Counter type for these types of data
 - A Histogram to **capture the latency of the HTTP requests** and the MySQL Queries
 - A Histogram **samples observations** (like latencies) and **counts them** in configurable **buckets**. It also provides a **sum of all observed values**.
 - The **default buckets** are intended to **cover a typical web/rpc request** from milliseconds to seconds

Monitoring a web app

- This is how I'm going to define the data types in Python for Prometheus:

```
from prometheus_client import Counter, Histogram

FLASK_REQUEST_LATENCY = Histogram('flask_request_latency_seconds', 'Flask Request Latency',
                                   ['method', 'endpoint'])
FLASK_REQUEST_COUNT = Counter('flask_request_count', 'Flask Request Count',
                              ['method', 'endpoint', 'http_status'])

MYSQL_REQUEST_LATENCY = Histogram('mysql_query_latency_seconds', 'MYSQL Query Latency',
                                   ['query'])
MYSQL_REQUEST_COUNT = Counter('mysql_query_count', 'Flask Request Count',
                              ['query'])
```

Monitoring a web app

- This is how we can calculate the latency of a query:

```
start_time = time.time()

sql = "select * from table"
# do the query

query_latency = time.time() - start_time

MYSQL_REQUEST_LATENCY.labels(sql[:50]).observe(query_latency)
MYSQL_REQUEST_COUNT.labels(sql[:50]).inc()
```


Demo

Monitor a web application with Prometheus

Demo

Monitor a web application with Prometheus - apdex score

Monitoring a web app

Prometheus with Java Spring boot

Monitoring a web app

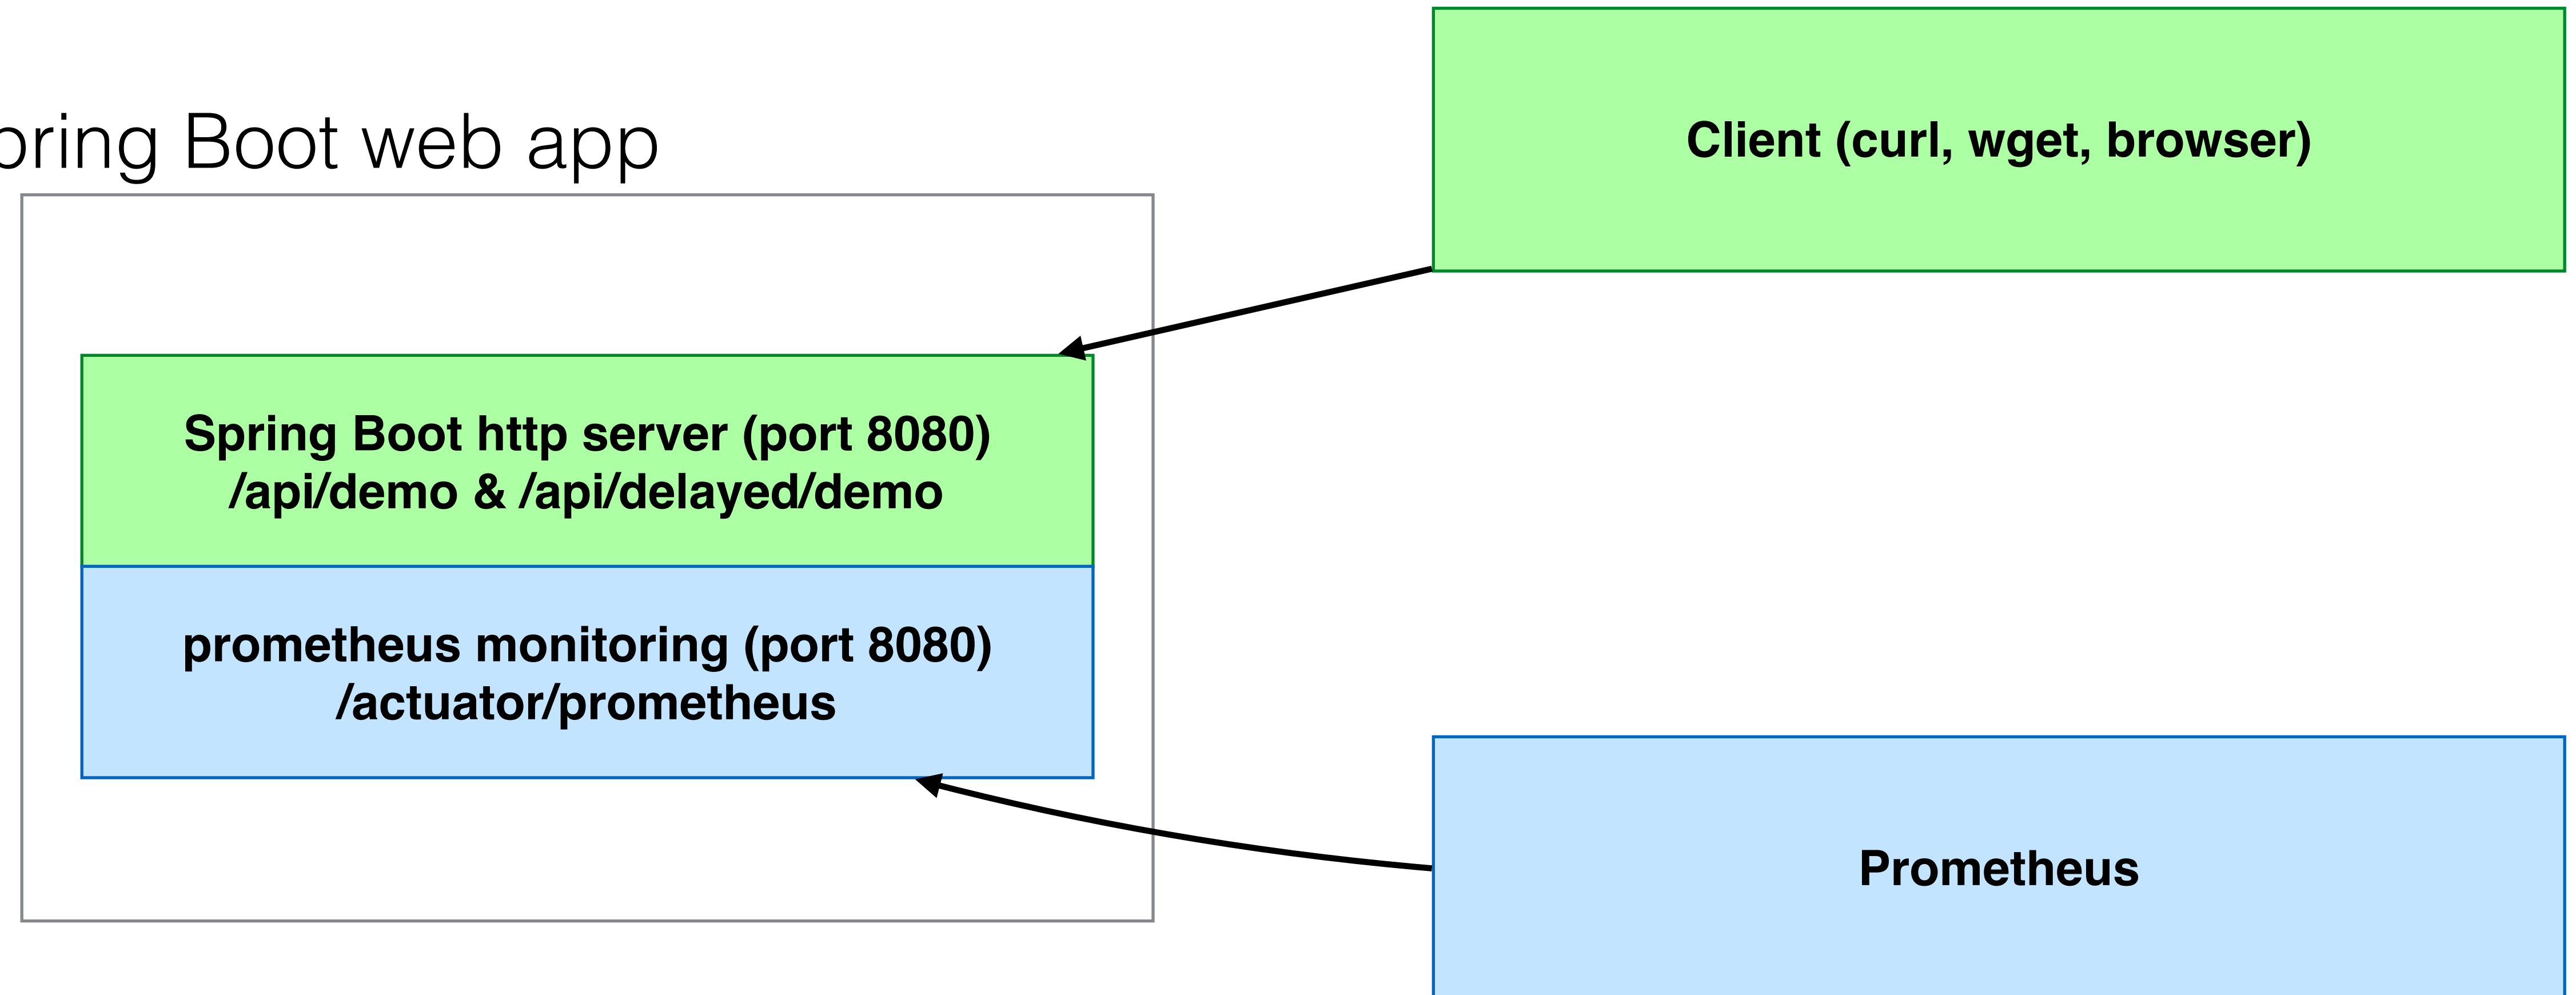
- Introduction
 - Application's **health + Metrics**
 - Notice **unwanted** behaviour
 - Monoliths as well as microservices
 - Crucial in microservices architecture
 - “To measure is to know”

Monitoring a web app

- We are going to integrate **prometheus monitoring** with a **web application** based on **Java Spring Boot**
 - We will use following :
 - **Spring Boot**
 - **Spring Boot Actuator**
 - **Micrometer**
 - We are also going to do a demo with an example.

The web app

Spring Boot web app



Monitoring a web app

- **Spring Boot Actuator**
 - Sub-Project of Spring boot
 - **Production** ready endpoints
 - **/actuator** is the common **prefix** of these endpoints
 - **Protected** by default
 - Adjustable in **application.properties**
 - Expose all: **management.endpoints.web.exposure.include=***

Monitoring a web app

- **Micrometer**
 - **Vendor-Neutral** application metrics facade
 - Support for **Prometheus** and many others:
 - AWS Cloudwatch, Datadog, InfluxDB/Telegraf, New Relic, ...
 - **Transforms** /actuator/metrics data into data your monitoring system understands
 - Only a **vendor-specific** micrometer dependency in your application is required

Monitoring a web app

- Micrometer
- pom.xml example

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>

<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-core</artifactId>
</dependency>

<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-registry-prometheus</artifactId>
</dependency>
```

Monitoring a web app

- Spring Boot
- Code example

```
...
import io.micrometer.core.instrument.Metrics;
...
private Counter runCounter = Metrics.counter("runCounter");
...
@GetMapping("/api/demo")
@Timed
public String apiUse() throws InterruptedException {
    runCounter.increment();
    log.info("Hello world app accessed on /api/demo");
    return "Hello world";
}
```

Demo

Monitor and instrument a Spring boot application

Grafana

Grafana Provisioning

Grafana Provisioning

- In one of the first lectures I showed you how to **setup Grafana using the UI**
- Rather than using the UI, you can also **use yaml and json files** to provision Grafana with datasources and dashboards
- This is a much more **powerful way** of using Grafana, as you can test new dashboards first on a **dev / test server**, then **import** the newly created dashboards to **production**
 - You can do the import manually through the UI, or **using yaml and json files**
 - When using files, you can keep files within **version control** to keep changes, revisions and backups

Grafana Provisioning

- The configuration of Grafana is all kept in `/etc/grafana/`:

`/etc/grafana/`:

```
-rw-r----- 1 root grafana 14K Jul 17 12:30 grafana.ini  
-rw-r----- 1 root grafana 3.4K Jul 17 12:30 ldap.toml  
drwxr-xr-x 4 root grafana 4.0K Jul 17 13:15 provisioning/
```

`/etc/grafana/provisioning/`:

```
drwxr-xr-x 2 root grafana 4.0K Jul 17 14:56 dashboards/  
drwxr-xr-x 2 root grafana 4.0K Jul 17 15:34 datasources/
```

- The data is kept in `/var/lib/grafana/`:

`/var/lib/grafana/`:

```
drwxr-xr-x 2 root root 4.0K Jul 17 15:47 dashboards/  
-rw-r----- 1 grafana grafana 500K Jul 17 15:48 grafana.db  
drwxr-x--- 2 grafana grafana 4.0K Jul 17 12:31 plugins/  
drwx----- 5 grafana grafana 4.0K Jul 17 12:40 sessions/
```

Grafana Provisioning

- You can change the database & paths in `/etc/grafana/grafana.ini`

```
[paths]
# Path to where grafana can store temp files, sessions, and the sqlite3 db (if that is used)
;data = /var/lib/grafana

# Directory where grafana can store logs
;logs = /var/log/grafana

# Directory where grafana will automatically scan and look for plugins
;plugins = /var/lib/grafana/plugins

# folder that contains provisioning config files that grafana will apply on startup and while running.
;provisioning = conf/provisioning
...
[database]
# Either "mysql", "postgres" or "sqlite3", it's your choice
;type = sqlite3
;host = 127.0.0.1:3306
;name = grafana
;user = root
# If the password contains # or ; you have to wrap it with triple quotes. Ex """"#password;""""
;password =
```

Demo

Grafana Provisioning

Prometheus Use Cases

Cloudwatch exporter

Use Cases - Cloudwatch

- Cloudwatch exporter
 - Installation
- Configuration (exporter + AWS)
- Charges + measuring them
- Querying metrics

Prometheus Use Cases

Consul integration

Consul integration

- Consul is a **distributed, highly available** solution providing:
 - A Service Mesh
 - Service Discovery
 - Health checks for your services
 - A Key-Value store
 - Secure Service Communications
 - Multi-datacenter support
- Consul is often deployed **in conjunction with Docker**

Consul integration

- There are **2 integrations** that are interesting to use:
 - 1) Prometheus can **scrape** Consul's metrics and provide you with all sorts of information about your running services
 - Consul provides **Service Discovery**, so it knows where services are running and what the **current state** of it is
 - 2) Consul can be integrated within Prometheus to **automatically add the services as targets**
 - Consul will discover your services, and these can then be automatically added to Prometheus as a target

Consul integration

- In the next demo I'll focus on the **Prometheus integration** with Consul, not really on implementing consul itself
 - I'll show you the **installation of consul**, but not how to integrate consul with your infrastructure (it's out of scope for this Prometheus course)
 - I'll **manually register a service to consul**, rather than setting up service discovery
 - In production environments, where you have a lot of services, service discovery using consul will allow you to register your services automatically
- If you are interested in Consul, have a look at the documentation at <https://www.consul.io/> and you can find the service registrator for docker at <https://github.com/gliderlabs/registrator>

Demo

Consul integration

Prometheus Use Cases

EC2 auto discovery

EC2 auto discovery

- **Service discovery** is the automatic detection of devices and services offered by these devices on a computer network.
- In this Use Case we will:
 - Create prerequisites in AWS (IAM role, Security Groups, EC2 instances)
 - Alter Prometheus config (*/etc/prometheus/prometheus.yml*)
 - Query the data in Grafana
 - <https://github.com/in4it/prometheus-course/blob/master/use-cases/ec2-auto-discovery/lab.txt>

Prometheus on Kubernetes

Getting Kubernetes metrics