Course Introduction

- Thank you for taking this Prometheus Course
 - Open-source monitoring solution & time series database
 - Was built by Soundcloud
 - Very active developer and user community
 - Now its a standalone open source project
 - Joined the Cloud Native Computing Foundation in 2016
 - Ideal for monitoring on premise as well as cloud workloads

Course Overview

Introduction	Monitoring	Alerting	Internals	Use cases
What is Prometheus	Client Libraries	Introduction	Storage	Grafana Provisioning
Installing Prometheus & Grafana	Pushing metrics	Setting up alerts	Security	Cloudwatch Integration
Concepts	Querying			
Configuration	Service Discovery			
Monitoring Nodes	Exporters			
Architecture				

Course Objectives

- To be able to use Prometheus
- To get familiar with the Prometheus ecosystem
- To setup a monitoring platform using
 - Prometheus
- To create alerts in Prometheus
- To be able to query Prometheus data

Who is Edward Viaene

- My name is Edward Viaene
- I am a consultant and trainer in Cloud and Big Data technologies
- I'm a big advocate of Agile and DevOps techniques
- I held various roles from banking to startups
- I have a background in Unix/Linux, Networks, Security, Risk, and distributed computing
- Nowadays I specialize in everything that has to do with Cloud and DevOps

Who is Jorn Jambers

- My name is Jorn Jambers
- I am a freelance DevOps consultant and trainer
- DevOps advocate
- Worked in banks, consultancy companies and startups
 - In the latter I found my passion for DevOps
- I have a background in Unix/Linux, Hadoop, DBA, Networks, automations
- Today I help companies succeed on the public cloud

Online Training

- Online training on Udemy
 - DevOps, Distributed Computing, Cloud, Big Data
 - Using online video lectures
 - 40,000+ enrolled students in 100+ countries

Introduction

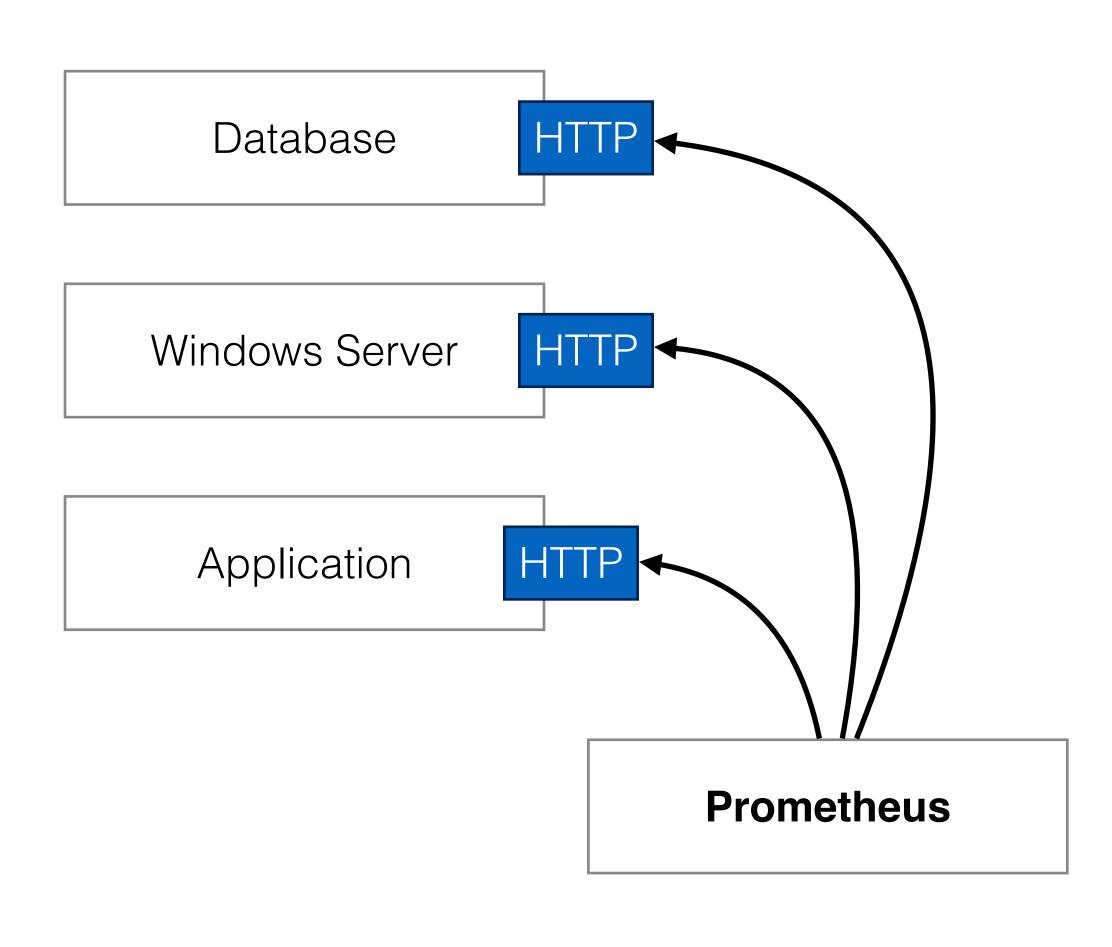
- Prometheus is an Open source monitoring solution
- Started at SoundCloud around 2012-2013, and was made public in early 2015
- Prometheus provides Metrics & Alerting
- It is inspired by Google's **Borgmon**, which uses time-series data as a datasource, to then send alerts based on this data
- It fits very well in the cloud native infrastructure
 - Prometheus is also a member of the CNCF (Cloud Native Foundation)

 In Prometheus we talk about **Dimensional Data**: time series are identified by metric name and a set of key/value pairs

Metric name	Label	Sample
Temperature	location=outside	90

- Prometheus includes a Flexible Query Language
- Visualizations can be shown using a built-in expression browser or with integrations like Grafana
- It stores metrics in memory and local disk in an own custom, efficient format
- It is written in **Go**
- Many client libraries and integrations available

How does Prometheus work?



- Prometheus collects metrics from monitored targets by scraping metrics HTTP endpoints
 - This is fundamentally different than other monitoring and alerting systems, (except this is also how Google's Borgmon works)
 - Rather than using custom scripts that check on particular services and systems, the monitoring data itself is used
- Scraping endpoints is much more efficient than other mechanisms, like 3rd party agents
 - A single prometheus server is able to ingest up to one million samples per second as several million time series

Installation

Prometheus Installation

- I will install Prometheus using scripts from our GitHub repository (https://github.com/
 in4it/prometheus-course)
 - They will work on any modern Linux distribution
- I'll install it on a DigitalOcean droplet
 - Feel free to use the scripts with any Cloud Provider, Virtual Machine, or Docker image, as long as it's a recent Linux distribution
 - To get a free \$100 coupon on DigitalOcean, valid for 60-days with a valid payment method added, use the following link:
 - https://m.do.co/c/b71b388ab76f
 - \$10 is enough to run a 2 **GB memory droplet** for one month

Prometheus Installation

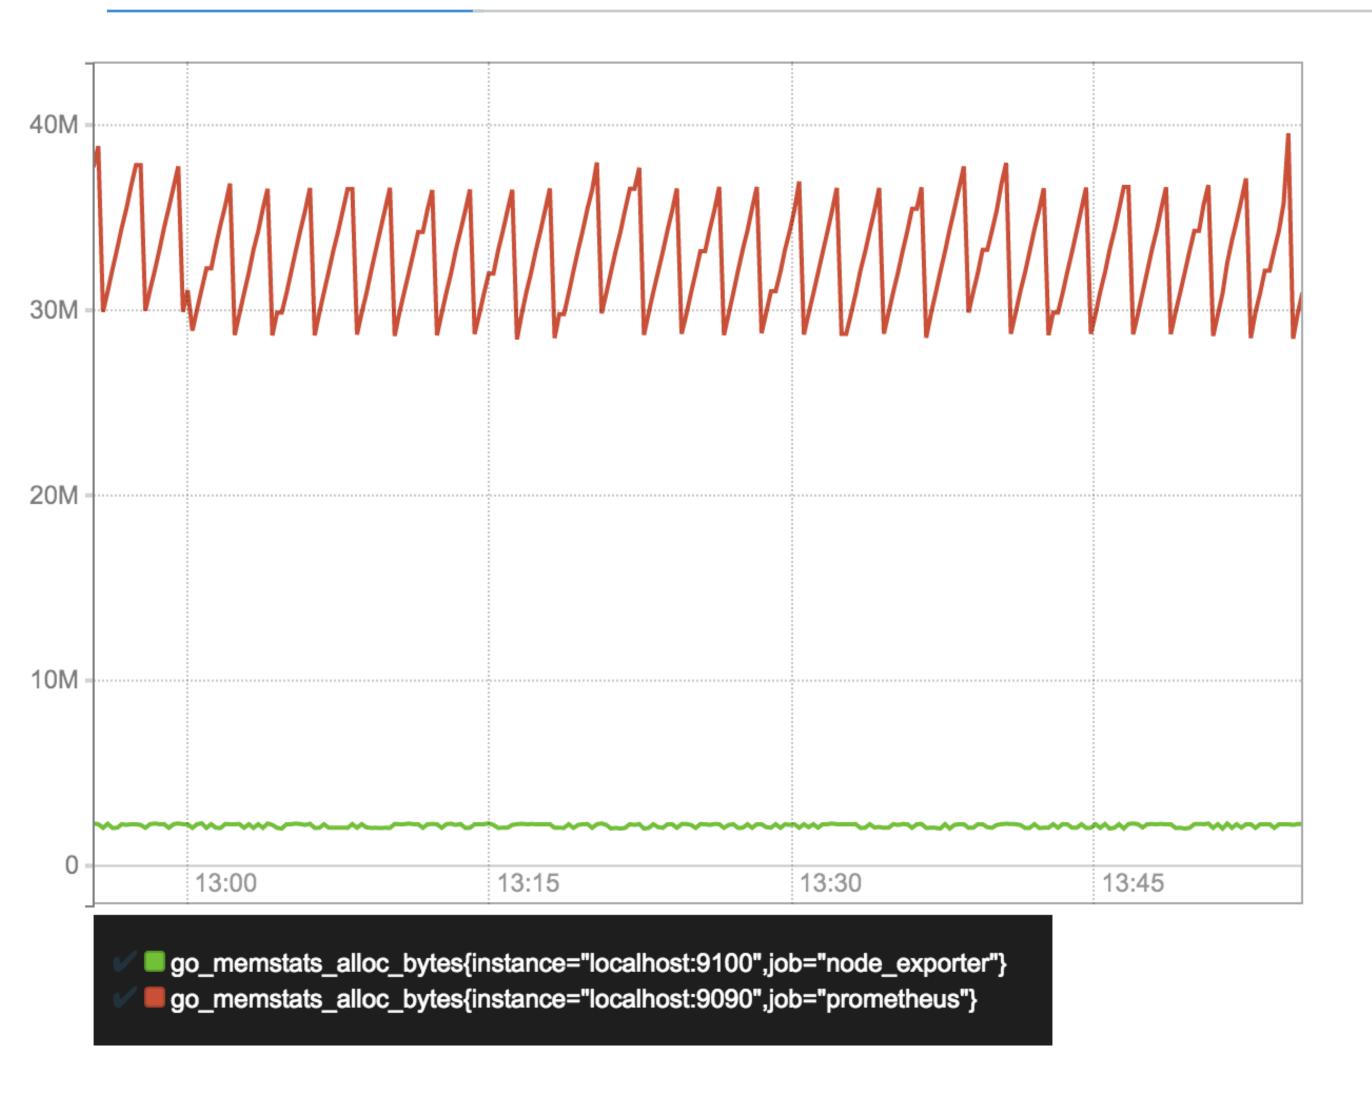
- If you do not want to use the provided scripts, you can download the full distribution from https://github.com/prometheus/prometheus/releases
- MacOS, Windows, Linux, and some Unix distributions are supported
- After extracting you'll get a **prometheus executable** (prometheus.exe for windows), which you can use to run prometheus, for example:
 - ./prometheus --config.file /path/to/prometheus.yaml
- It's best to use the scripts we provided so that your environment is the same as ours when you follow the demos

Demo

Installing Prometheus & Grafana

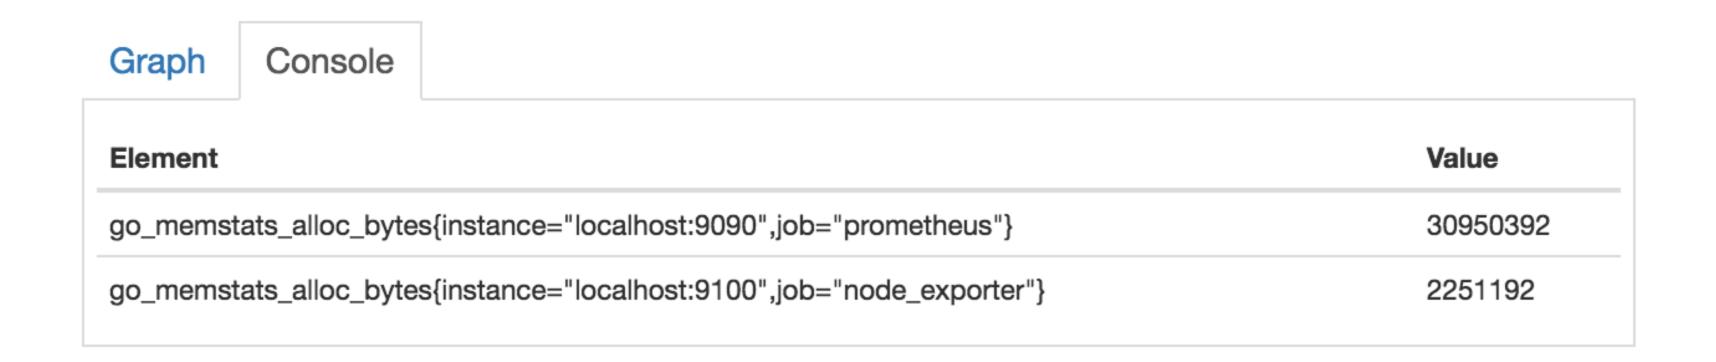
Basic Concepts

Concepts



- All data is stored as time series
 - Every time series is identified by the "metric name" and a set of key-value pairs, called labels
 - metric: go_memstat_alloc_bytes
 - instance=localhost:9090
 - job=prometheus

- The time series data also consists of the actual data, called Samples:
 - It can be a float64 value
 - or a millisecond-precision timestamp



- The notation of time series is often using this notation:
 - <metric name>{<label name>=<label value>, ...}
 - For example:
 - node_boot_time{instance="localhost:9100",job="node_exporter"}

Prometheus Configuration file

Prometheus Configuration

- The configuration is stored in the Prometheus configuration file, in yaml format
 - The configuration file can be changed and applied, without having to restart Prometheus
 - A reload can be done by executing kill -SIGHUP <pid>
- You can also pass parameters (flags) at startup time to ./prometheus
 - Those parameters cannot be changed without restarting Prometheus
- The configuration file is passed using the flag --config.file

Prometheus Configuration

The default configuration looks like this:

```
# my global config
global:
                      15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
  scrape_interval:
  evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.
  # scrape_timeout is set to the global default (10s).
# Alertmanager configuration
alerting:
  alertmanagers:
  - static_configs:
    - targets:
      # - alertmanager:9093
# Load rules once and periodically evaluate them according to the global 'evaluation_interval'.
rule_files:
 # - "first_rules.yml"
  # - "second_rules.yml"
```

Prometheus Configuration

- To scrape metrics, you need to add configuration to the prometheus config file
- For example, to scrape metrics from prometheus itself, the following code block is added by default

Demo

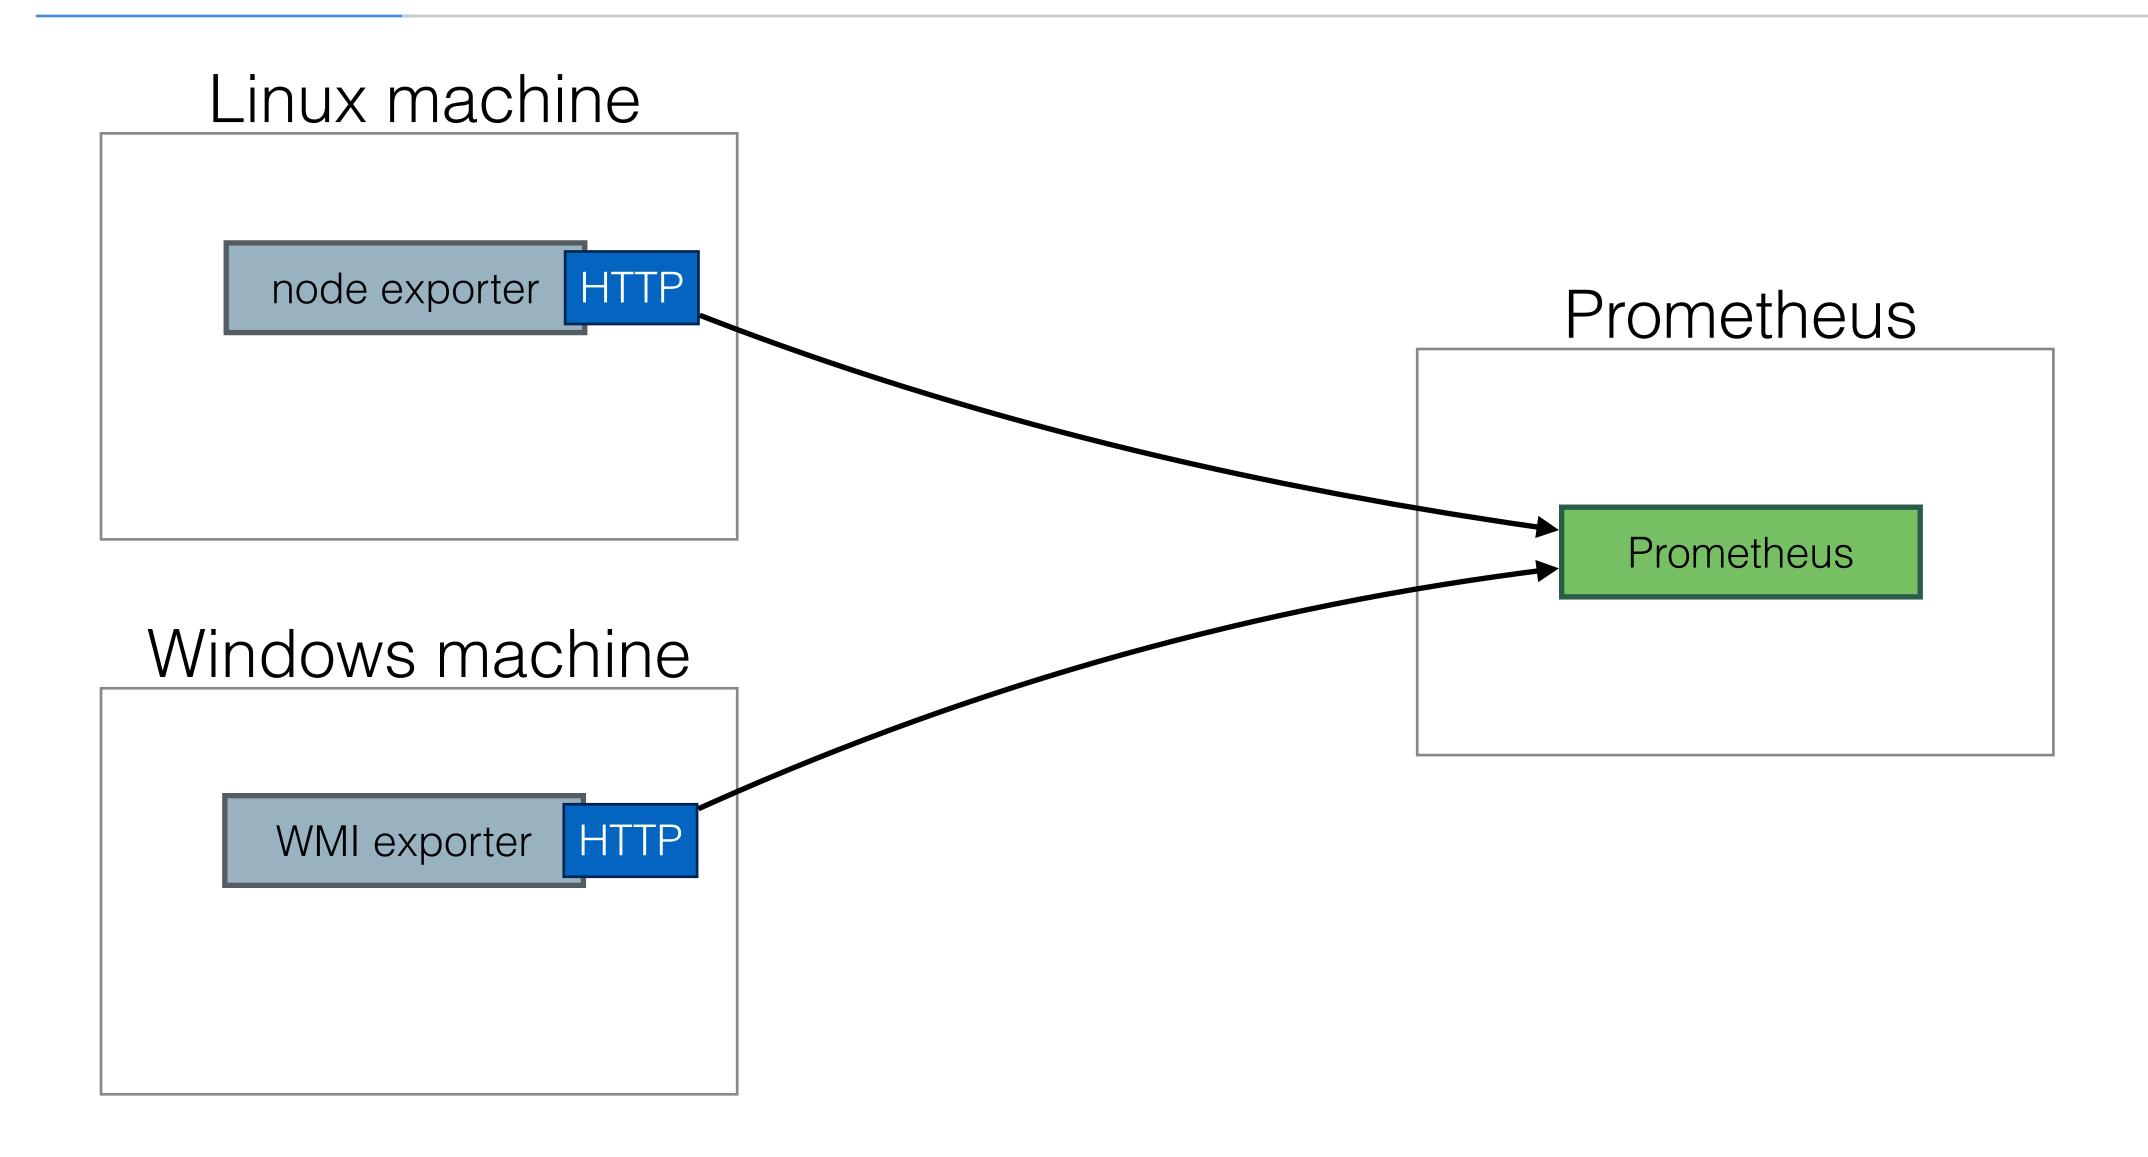
Prometheus Configuration

Monitoring Nodes

Monitor nodes

- To monitor nodes, you need to install the node-exporter
- The node exporter will expose machine metrics of Linux / *Nix machines
 - For example: cpu usage, memory usage
- The node exporter can be used to monitor machines, and later on, you can create alerts based on these ingested metrics
- For Windows, there's a WMI exporter (see https://github.com/martinlindhe/wmi_exporter)

Monitor nodes



Demo

Node exporter

Demo

WMI exporter

Monitoring

Monitoring - Introduction

- Client Libraries
- Pushing Metrics
- Querying
- Service Discovery
- Exporters

Client Libraries

Client Libraries - Introduction

- Instrumenting your code
- Libraries
 - Official: Go, Java/Scala, Python, Ruby
 - Unofficial: Bash, C++, Common Lisp, Elixir, Erlang, Haskell, Lua for Nginx, Lua for Tarantool, .NET / C#, Node.js, PHP, Rust
- No client library available?
 - Implement it yourself in one of the supported exposition formats

Client Libraries - Introduction

- Exposition formats:
 - Simple text-based format
 - Protocol-buffer format (Prometheus 2.0 removed support for the protocol-buffer format)

```
metric_name [
   "{" label_name "=" `"` label_value `"` { "," label_name "=" `"` label_value `"` } [ "," ] "}"
] value [ timestamp ]

node_filesystem_avail_bytes{device="/dev/vda1",fstype="ext4",mountpoint="/"} 4.9386491904e+10
node_filesystem_avail_bytes{device="/dev/vda15",fstype="vfat",mountpoint="/boot/efi"} 1.05903104e+08
node_filesystem_avail_bytes{device="lxcfs",fstype="fuse.lxcfs",mountpoint="/var/lib/lxcfs"} 0
node_filesystem_avail_bytes{device="tmpfs",fstype="tmpfs",mountpoint="/run"} 2.01273344e+08
node_filesystem_avail_bytes{device="tmpfs",fstype="tmpfs",mountpoint="/run/lock"} 5.24288e+06
```

Client Libraries - Introduction

- 4 types of metrics
 - Counter
 - Gauge
 - Histogram
 - Summary

Counter

A value that only goes up (e.g. Visits to a website)

Gauge

Single numeric value that can go up and down (e.g. CPU load, temperature)

Histogram

Samples observations (e.g. request durations or response sizes) and these observations get counted into **buckets**. Includes (_count and _sum)

Main purpose is calculating quantiles

Summary

Similar to a *histogram*, a *summary* samples observations (e.g. request durations or response sizes). A summary also provides a total count of observations and a sum of all observed values, it calculates configurable quantiles over a sliding time window.

Example: You need 2 counters for calculating the latency

- 1) total request(_count)
- 2) the total latency of those requests (**_sum**)

Take the rate() and divide = average latency

Prometheus

Instrumentation- Python

- https://github.com/prometheus/client_python
- Officially supported language
- pip install prometheus_client
- Supported metrics: Counter, Gauge, Summary and Histogram

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/<name>", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/jorn", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
import random, time
from flask import Flask, render_template_string, abort
from prometheus_client import generate_latest, REGISTRY, Counter, Gauge, Histogram
app = Flask(__name___)
REQUESTS = Counter('http_requests_total', 'Total HTTP Requests (count)', ['method', 'endpoint', 'status_code'])
IN_PROGRESS = Gauge('http_requests_inprogress', 'Number of in progress HTTP requests')
TIMINGS = Histogram('http_request_duration_seconds', 'HTTP request latency (seconds)')
@app.route('/')
@TIMINGS.time()
@IN_PROGRESS.track_inprogress()
def hello_world():
    REQUESTS.labels(method='GET', endpoint="/", status_code=200).inc() # Increment the counter
    return 'Hello, World!'
@app.route('/prometheus-course/<name>')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def index(name):
    REQUESTS.labels(method='GET', endpoint="/prometheus-course/jorn", status_code=200).inc()
    return render_template_string('<b>Hello {{name}} welcome!</b>!', name=name)
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
```

```
@app.route('/metrics')
@IN_PROGRESS.track_inprogress()
@TIMINGS.time()
def metrics():
    REQUESTS.labels(method='GET', endpoint="/metrics", status_code=200).inc()
    return generate_latest(REGISTRY)

if __name__ == "__main__":
    app.run(host='0.0.0.0')
```

Prometheus

Instrumentation - Go

- https://github.com/prometheus/client_golang
- Officially supported language
- Easy to implement:

```
package main

import (
    "github.com/prometheus/client_golang/prometheus/promhttp"
    "net/http"
)

func main() {
    http.Handle("/metrics", promhttp.Handler())
    panic(http.ListenAndServe(":8080", nil))
}
```

• Supported metrics: Counter, Gauge, Summary and Histogram

Gauge

```
import "github.com/prometheus/client_golang/prometheus"
var jobsInQueue = prometheus.NewGauge(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs queued",
   },
func init(){
    promtheus.MustRegister(jobsQueued)
func enqueueJob(job Job) {
    queue.Add(job)
    jobsInQueue.Inc()
func runNextJob() {
    job := queue.Dequeue()
   jobsInQueue.Dec()
   job.Run()
```

Gauge

```
import "github.com/prometheus/client_golang/prometheus"
var jobsQueued = prometheus.NewGauge(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs queued",
func init(){
    promtheus.MustRegister(jobsQueued)
func enqueueJob(job Job) {
    queue.Add(job)
    jobsQueued.Inc()
func runNextJob() {
    job := queue.Dequeue()
    jobsInQueue.Dec()
   job.Run()
```

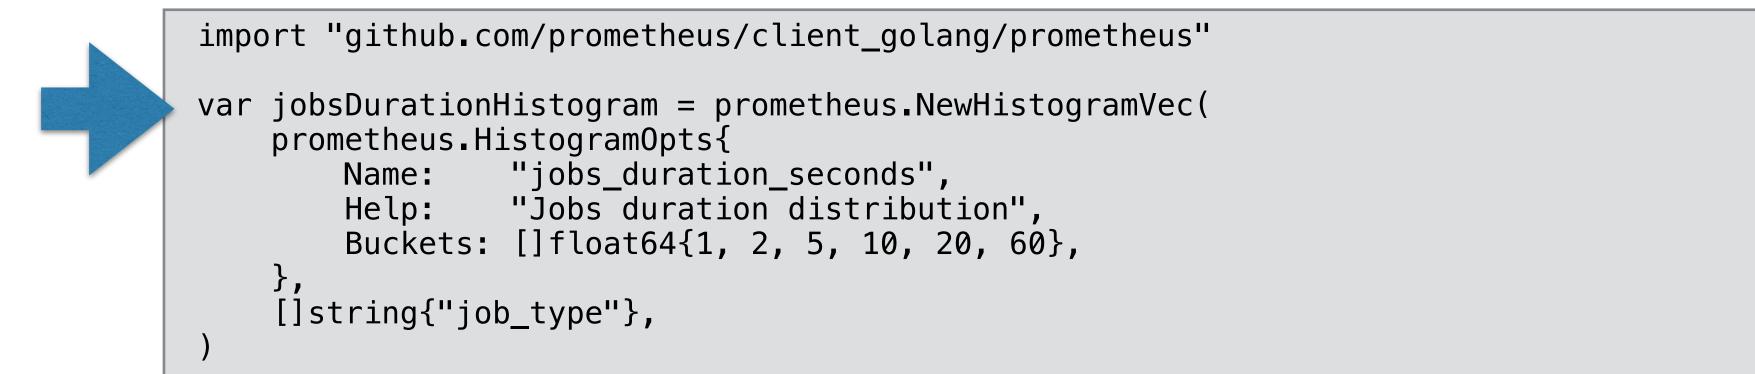
Adding labels

```
import "github.com/prometheus/client_golang/prometheus"
var jobsQueued = prometheus.NewGaugeVec(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs in the queue",
    []string{"job_type"},
func init(){
    promtheus.MustRegister(jobsQueued)
func enqueueJob(job Job) {
    queue.Add(job)
    jobsInQueue.WithLabelValues(job.Type()).Inc()
func runNextJob() {
    job := queue.Dequeue()
    jobsInQueue.WithLabelValues(job.Type()).Dec()
    job.Run()
```

Adding labels

```
import "github.com/prometheus/client_golang/prometheus"
var jobsQueued = prometheus.NewGaugeVec(
    prometheus.GaugeOpts{
        Name: "jobs_queued",
        Help: "Current number of jobs in the queue",
    []string{"job_type"},
func init(){
    promtheus.MustRegister(jobsQueued)
func enqueueJob(job Job) {
    queue.Add(job)
    jobsQueued.WithLabelValues(job.Type()).Inc()
func runNextJob() {
    job := queue.Dequeue()
    jobsQueued.WithLabelValues(job.Type()).Dec()
    job.Run()
```

Histogram



jobsDurationHistogram.WithLabelValues(job.Type()).Observe(duration.Seconds())

start := time.Now()

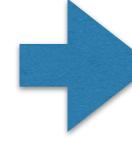
duration := time.Since(start)

job.Run()

Histogram

```
import "github.com/prometheus/client_golang/prometheus"
var jobsDurationHistogram = prometheus.NewHistogramVec(
    prometheus.HistogramOpts{
              "jobs_duration_seconds",
       Name:
              "Jobs duration distribution",
        Buckets: []float64{1, 2, 5, 10, 20, 60},
    []string{"job_type"},
start := time.Now()
job.Run()
duration := time.Since(start)
jobsDurationHistogram.WithLabelValues(job.Type()).Observe(duration.Seconds())
```

Summary



prometheus.NewSummary()

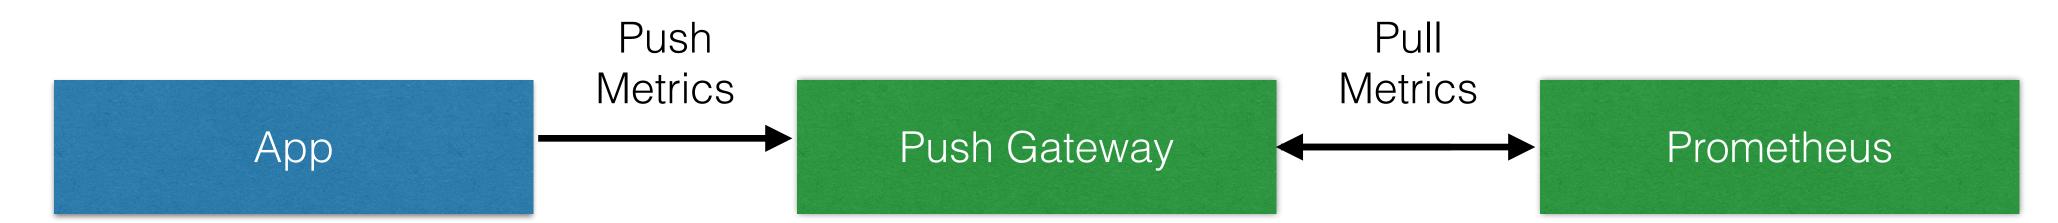
Prometheus

Pushing metrics

Pushing Metrics - Introduction

https://github.com/prometheus/pushgateway

Diagram



Pushing Metrics - Introduction

- Sometimes metrics cannot be scraped
 Example: batch jobs, servers are not reachable due to NAT, firewall
- Pushgateway is used as an intermediary service which allows you to push metrics.
- · Pitfalls
 - Most of the times this is a single instance so this results in a SPOF
 - Prometheus's automatic instance health monitoring is not possible
 - The Pushgateway never forgets the metrics unless they are deleted via the api example:

curl -X DELETE http://localhost:9091/metrics/job/prom_course/instance/localhost

Pushing Metrics - Introduction

- Only 1 valid use case for the Pushgateway
 - Service-level batch jobs and not related to a specific machine
- If NAT or/both firewall is blocking you from using the pull mechanism
 - Move the Prometheus server on the same network

Pushing Metrics - Python Example

```
from prometheus_client import CollectorRegistry, Gauge, push_to_gateway

registry = CollectorRegistry()
g = Gauge('job_last_success_unixtime', 'Last time the course batch job has finished', registry=registry)
g.set_to_current_time()
push_to_gateway('localhost:9091', job='batchA', registry=registry)
```

- Pushgateway functions take a grouping key.
 - push_to_gateway replaces metrics with the same grouping key
 - pushadd_to_gateway only replaces metrics with the same name and grouping key
 - delete_from_gateway deletes metrics with the given job and grouping key.

Prometheus

Pushing Metrics - Go

Pushing Metrics - Go Example

Go example:

```
package main
import (
   "flag"
    "log"
   "net/http"
   "github.com/prometheus/client_golang/prometheus/promhttp"
    "github.com/prometheus/client_golang/prometheus/push"
gatewayUrl:="http://localhost:9091/"
throughputGuage := prometheus.NewGauge(prometheus.GaugeOpts{
       Name: "throughput",
        Help: "Throughput in Mbps",
throughputGuage Set (800)
if err := push.Collectors(
        "throughput_job", push.HostnameGroupingKey(),
       gatewayUrl, throughputGuage
); err != nil {
    fmt.Println("Could not push completion time to Pushgateway:", err)
```

Prometheus

Querying

Querying Metrics - Introduction

- Prometheus provides a functional expression language called PromQL
 - Provides built in operators and functions
 - Vector-based calculations like Excel
 - Expressions over time-series vectors
- PromQL is read-only
- Example:

```
100 - (avg by (instance) (irate(node_cpu_seconds_total{job='node_exporter',mode="idle"}[5m])) * 100)
```

Prometheus

Querying - Expressions

Querying Metrics - Introduction

• Instant vector - a set of time series containing a single sample for each time series, all sharing the same timestamp

Example: node_cpu_seconds_total

 Range vector - a set of time series containing a range of data points over time for each time series

Example: node_cpu_seconds_total[5m]

- **Scalar** a simple numeric floating point value Example: -3.14
- **String** a simple string value; currently unused Example: foobar

Prometheus

Querying - Operators

Querying Metrics - Introduction

Arithmetic binary operators

Example: - (subtraction), * (multiplication), / (division), % (modulo), ^ (power/exponentiation)

Comparison binary operators

Example: == (equal), != (not-equal), > (greater-than), < (less-than) ,>= (greater-or-equal), <= (less-or-equal)

Logical/set binary operators

Example: and (intersection), or (union), unless (complement)

Aggregation operators

Example:**sum** (calculate sum over dimensions), **min** (select minimum over dimensions) ,**max** (select maximum over dimensions), **avg** (calculate the average over dimensions), **stddev** (calculate population standard deviation over dimensions), **stdvar** (calculate population standard variance over dimensions), **count** (count number of elements in the vector), **count_values** (count number of elements with the same value), **bottomk** (smallest k elements by sample value), **topk** (largest k elements by sample value), **quantile** (calculate ϕ -quantile ($0 \le \phi \le 1$) over dimensions)

Demo

Querying

Service Discovery

- Definition:
 - **Service discovery** is the automatic detection of devices and services offered by these devices on a computer network.
- Not really a service discovery mechanism

```
static_configs:
    - targets: ['localhost:9090']
```

- Cloud support for (AWS, Azure, Google,...)
- Cluster managers (Kubernetes, Marathon, ...)
- Generic mechanisms (DNS, Consul, Zookeeper, ...)

Service Discovery - Example AWS

• EC2 Example: Add following config to /etc/prometheus/prometheus.yml

```
global:
    scrape_interval: 1s
    evaluation_interval: 1s

scrape_configs:
    - job_name: 'node'
    ec2_sd_configs:
    - region: eu-west-1
    access_key: PUT_THE_ACCESS_KEY_HERE
    secret_key: PUT_THE_SECRET_KEY_HERE
    port: 9100
```

- Make sure the user has the following IAM role: AmazonEC2ReadOnlyAccess
- Make sure you security groups allow access to port (9100, 9090)

EC2 Example:
 Only monitor instances started with the name PROD

```
global:
 scrape_interval: 1s
 evaluation_interval: 1s
scrape_configs:
  - job_name: 'node'
    ec2_sd_configs:
      - region: eu-west-1
        access_key: PUT_THE_ACCESS_KEY_HERE
        secret_key: PUT_THE_SECRET_KEY_HERE
        port: 9100
    relabel_configs:
        # Only monitor instances with a tag Name starting with "PROD"
      - source_labels: [__meta_ec2_tag_Name]
        regex: PROD.*
        action: keep
       # Use the instance ID as the instance label
      - source_labels: [__meta_ec2_instance_id]
        target_label: instance
```

• EC2 Example:

Relabel ip adress to instance id for convenience

```
global:
 scrape_interval: 1s
 evaluation_interval: 1s
scrape_configs:
  - job_name: 'node'
    ec2_sd_configs:
      - region: eu-west-1
        access_key: PUT_THE_ACCESS_KEY_HERE
        secret_key: PUT_THE_SECRET_KEY_HERE
        port: 9100
    relabel_configs:
        # Only monitor instances with a tag Name starting with "PROD"
      - source_labels: [__meta_ec2_tag_Name]
        regex: PROD.*
        action: keep
       # Use the instance ID as the instance label
      - source_labels: [__meta_ec2_instance_id]
        target_label: instance
```

Service Discovery - Example Kubernetes

Kubernetes Example:
 Add following config to /etc/prometheus/prometheus.yml

```
- job_name:'kubernetes'
   kubernetes_sd_configs:
       api servers:
         - https://kubernetes.default.svc
       in_cluster: true
       basic_auth:
       username: prometheus
       password: secret
       retry_interval:5s
- job_name:'kubernetes-service-endpoints'
   kubernetes_sd_configs:
       api_servers:
         - https://kube-master.prometheuscourse.com
       in_cluster: true
```

Service Discovery - Example DNS

Service Discovery - DNS

• DNS Example: Add following config to /etc/prometheus/prometheus.yml

Service Discovery - Example using file

Service Discovery - Using file

• File Example:

Add following config to /etc/prometheus/prometheus.yml

```
scrape_configs:
    - job_name: 'dummy' # This is a default value, it is
mandatory.
    file_sd_configs:
    - files:
    - targets.json
```

Format target.json

Exporters

Exporters - Introduction

- Build for exporting prometheus metrics from existing third-party metrics
- When Prometheus is not able to pull metrics directly(Linux sys stats, haproxy, ...)
- Examples:

MySQL server exporter

Memcached exporter

Consul exporter

Node/system metrics exporter

MongoDB

Redis

Many more....

https://prometheus.io/docs/instrumenting/exporters/

Exporters - Introduction

We are already using one :-)
 Check /etc/prometheus/prometheus.yml

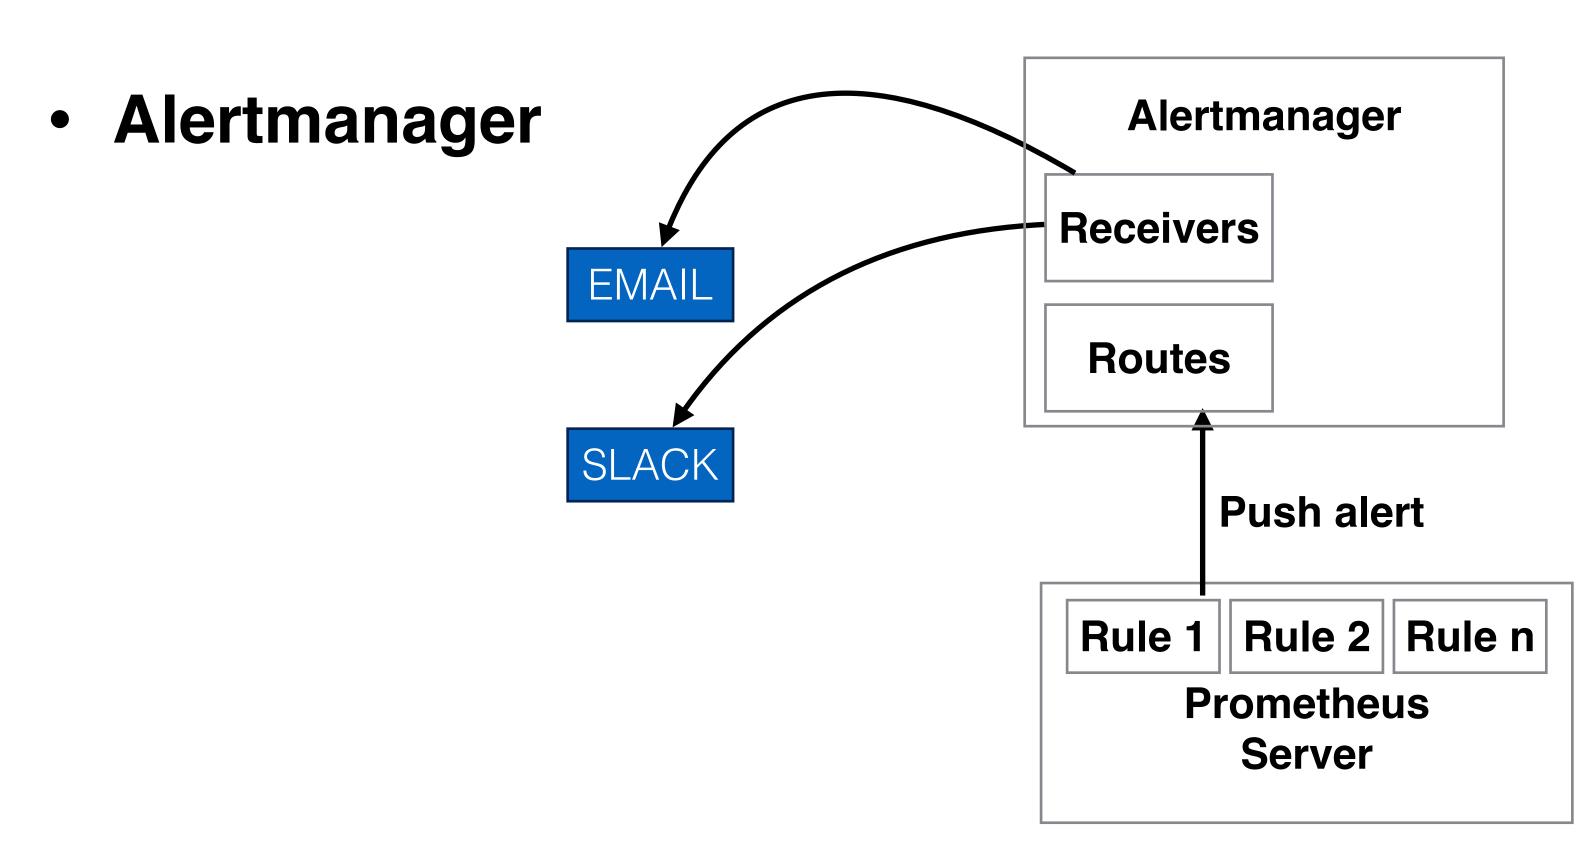
```
- job_name: 'node_exporter'
    scrape_interval: 5s
    static_configs:
    - targets: ['localhost:9100']
```

Alerting

Alerting - Introduction

Alerting - Introduction

- Alerting in Prometheus is separated into 2 parts
 - Alerting rules in Prometheus server



Alerting - Alerting rules

Alerting Rules

- Rules live in Prometheus server config
- Best practice to separate the alerts from the prometheus config
 - Add an include to /etc/prometheus/prometheus.yml
 rule_files:
 "/etc/prometheus/alert.rules"
- Alert format:

```
ALERT <alert name>
    IF <expression>
    [ FOR <duration> ]
    [ LABELS <label set> ]
    [ ANNOTATIONS <label set> ]
```

Alert example:

```
groups:
- name: example
rules:
- alert: cpuUsge
expr: 100 - (avg by (instance) (irate(node_cpu_seconds_total{job='node_exporter',mode="idle"}[5m])) * 100) >
95
for: 1m
labels:
severity: critical
annotations:
summary: Machine under healvy load
```

Alerting Rules

- Alerting rules allow you to define the alert conditions
- Alerting rules sent the alerts being fired to an external service
- The format of these alerts is in the Prometheus expression language

• Example:

```
groups:
- name: Important instance
rules:

# Alert for any instance that is unreachable for >5 minutes.
- alert: InstanceDown
    expr: up == 0
    for: 5m
    labels:
        severity: page
    annotations:
        summary: "Instance {{ $labels.instance }} down"
        description: "{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 5 minutes."
```

Alerting - Alertmanager

- Alertmanager handles the alerts fired by the prometheus server
- Handles deduplication, grouping and routing of alerts
- Routes alerts to receivers (Pagerduty, Opsgenie, email, Slack,...)

• Alertmanager Configuration (/etc/alertmanager/alertmanager.yml):

```
global:
 smtp_smarthost: 'localhost:25'
 smtp_from: 'alertmanager@prometheus.com'
 smtp_auth_username: ''
 smtp_auth_password: ''
templates:
- '/etc/alertmanager/template/*.tmpl'
route:
 repeat_interval: 1h
 receiver: operations-team
receivers:
- name: 'operations-team'
 email_configs:
 - to: 'operations-team+alerts@example.org'
 slack_configs:
 channel: '#prometheus-course'
   send_resolved: true
```

• Prometheus Configuration (/etc/prometheus/prometheus.yml):

```
# my global config
global:
                    15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
  scrape_interval:
  evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.
  # scrape_timeout is set to the global default (10s).
# Alertmanager configuration
alerting:
  alertmanagers:
  - static_configs:
    - targets:
       - localhost:9093
# Load rules once and periodically evaluate them according to the global 'evaluation_interval'.
rule_files:
 # - "first_rules.yml"
  # - "second_rules.yml"
# A scrape configuration containing exactly one endpoint to scrape:
# Here it's Prometheus itself.
scrape_configs:
 # The job name is added as a label `job=<job_name>` to any timeseries scraped from this config.
  - job_name: 'prometheus'
    # metrics_path defaults to '/metrics'
    # scheme defaults to 'http'.
    static_configs:
      - targets: ['localhost:9090']
```

- Concepts:
 - Grouping: Groups similar alerts into 1 notification
 - Inhibition: Silence other alerts if one specified alert is already fired
 - Silences: A simple way to mute certain notifications

High availability


- You can create a high available Alertmanager cluster using mesh config
- Do not load balance this service!
 - Use a list of Alertmanager nodes in Prometheus config
- All alerts are sent to all known Alertmanager nodes
- No need to monitor the monitoring
- Guarantees the notification is at least send once

Alert states:

Inactive - No rule is met

Pending - Rule is met but can be suppressed due to validations Firing - Alert is sent to the configured channel(mail, Slack,...)

• Runs on port :9093



Notifying multiple destinations

Setting up alerts

Setting up alerts - Demo

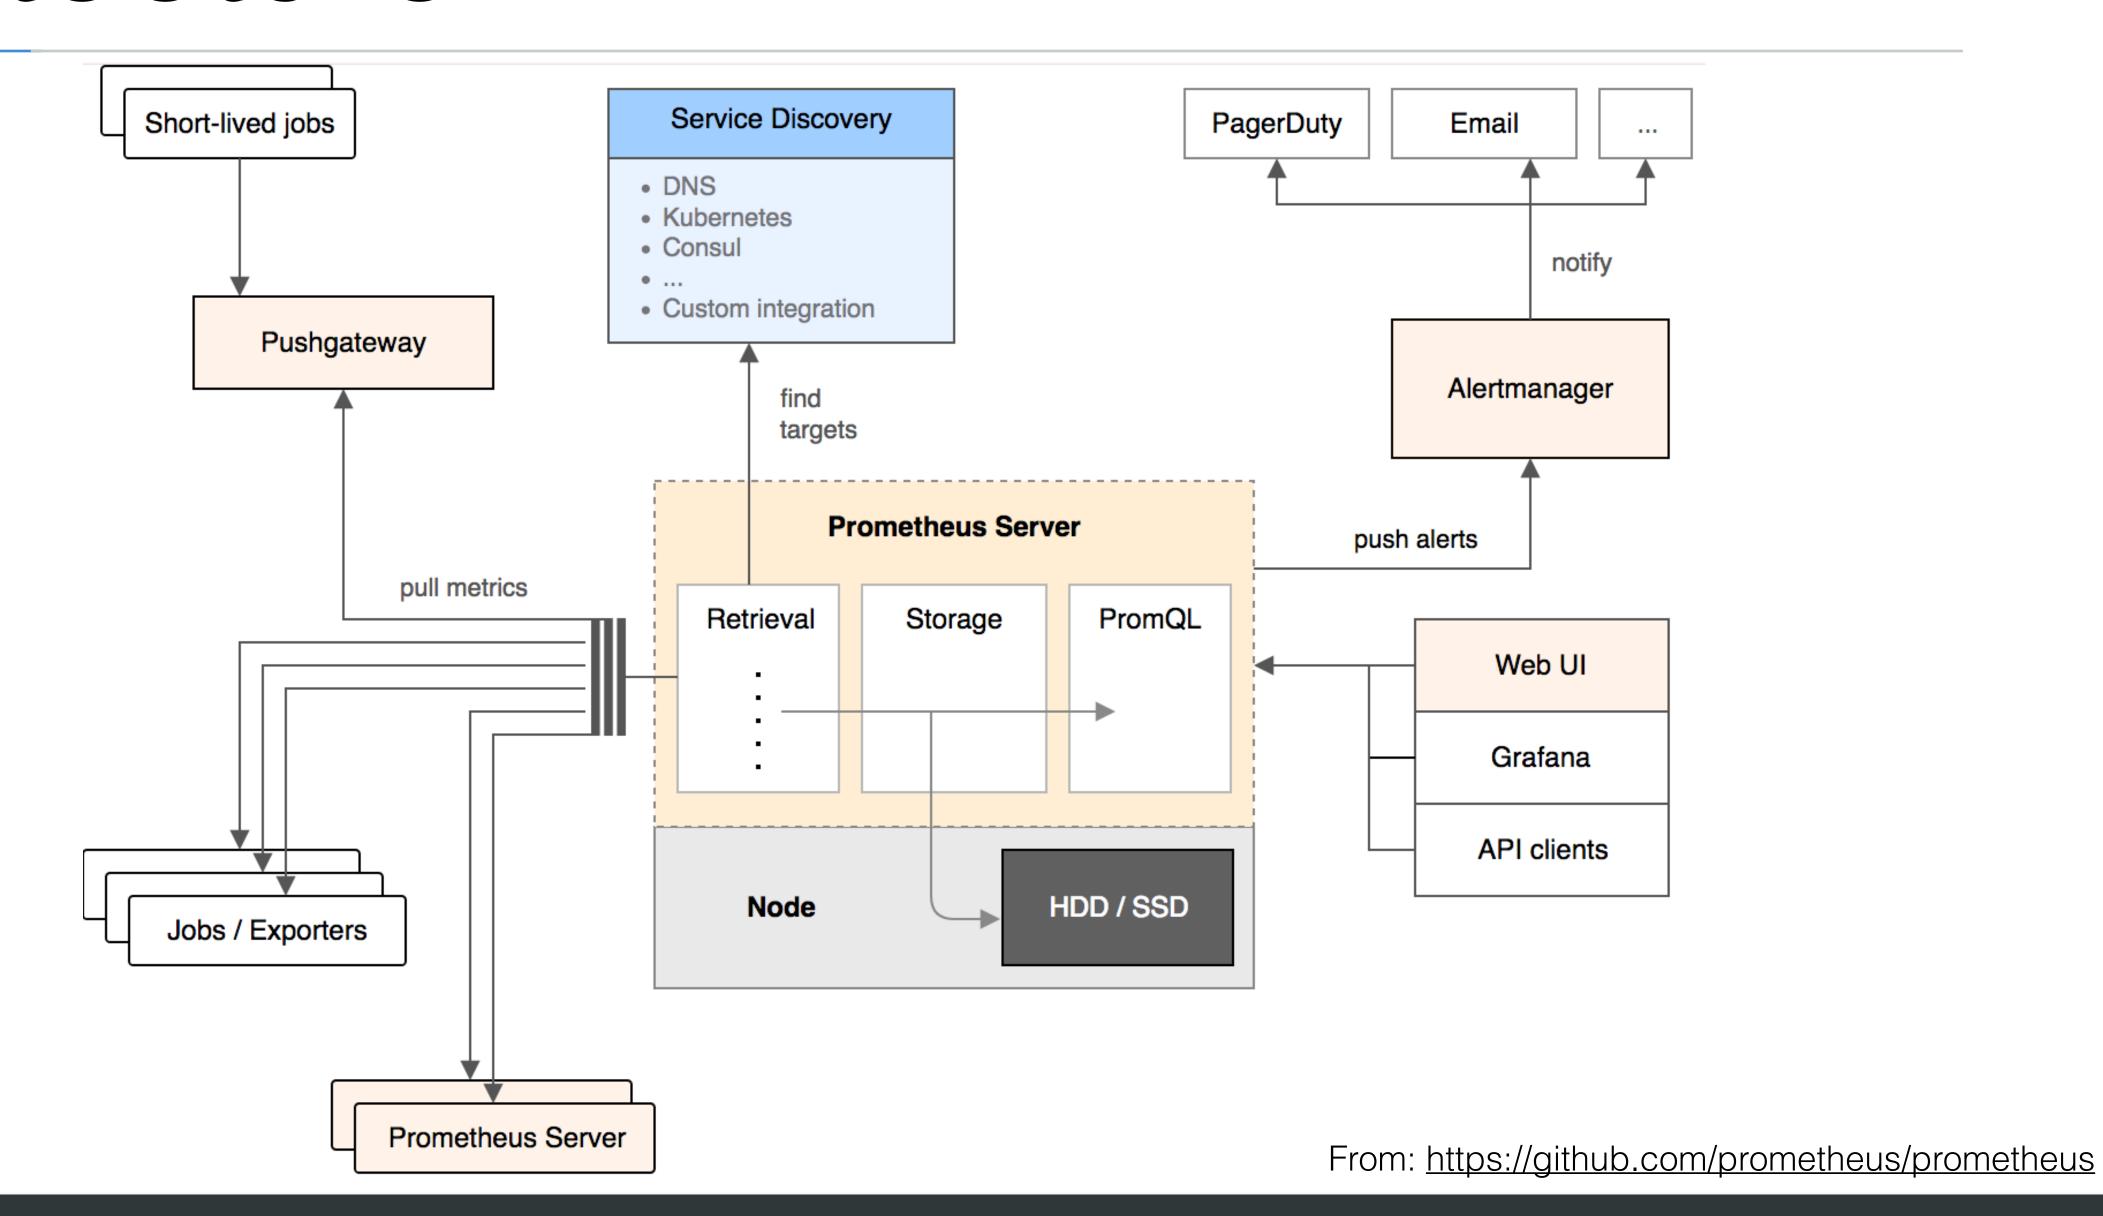
Setting up alerts

- Install Alertmanager
- Create config for the Alertmanager
 - Mail
 - Slack
- Alter prometheus config
- Setup an alert
 - See the notification coming in when an alert is fired

Prometheus internals

Architecture

Architecture



Prometheus

Storage

Storage

- You can use the default local on-disk storage, or optionally the remote storage system
 - Local storage: a local time series database in a custom Prometheus format
 - Remote storage: you can read/write samples to a remote system in a standardized format
 - Currently it uses a snappy-compressed protocol buffer encoding over HTTP, but might change in the future (to use gRPC or HTTP/2)

Remote Storage

- Remote storage is primarily focussed at long term storage
- Currently there are adapters available for the following solutions:

AppOptics: write	Graphite: write
Chronix: write	InfluxDB: read and write
Cortex: read and write	OpenTSDB: write
CrateDB: read and write	PostgreSQL/TimescaleDB: read and write
Gnocchi: write	SignalFx: write

Source: https://prometheus.io/docs/operating/integrations/#remote-endpoints-and-storage

- Prometheus >=2.0 uses a **new storage engine** which dramatically increases scalability
- Ingested samples are grouped in blocks of two hours
- Those 2h samples are stored in **separate directories** (in the data directory of prometheus)
- Writes are batched and written to disk in chunks, containing multiple data points

directory 1

directory 2

directory 3

2h of data:

chunks/000001 chunks/000002 2h of data:

chunks/000001

2h of data:

chunks/000001 chunks/000002

- Every directory also has an index file (index) and a metadata file (meta.json)
- It stores the **metric names** and the **labels**, and **provides an index** from the metric names and labels to the series in the chunk files

directory 1

chunks/000001 chunks/000002 meta.json index directory 2

chunks/000001
meta.json
index

directory 3

chunks/000001 chunks/000002 **meta.json index**

- The most recent data is kept in memory
- You don't want to loose the in-memory data during a crash, so the data also needs to be persisted to disk. This is done using a write-ahead-log (WAL)

directory 1

chunks/000001 chunks/000002 meta.json index directory 2

chunks/000001 meta.json index directory 3

chunks/000001 chunks/000002 meta.json index wal:

000001

- Write Ahead Log (WAL)
 - It's quicker to append to a file (like a log) than making (multiple) random read/writes
 - If there's a server crash and the data from memory is lost, then the WAL will be replayed
 - This way, no data will be lost or corrupted during a crash

- When series gets deleted, a tombstone file gets created
- This is more efficient than immediately deleting the data from the chunk files, as the actual delete can happen at a later time (e.g. when there's not a lot of load)

directory 1

2h of data:

chunks/000001 chunks/000002 meta.json index **tombstone** directory 2

2h of data:

chunks/000001 chunks/000002 meta.json index **tombstone** directory 3

2h of data:

chunks/000001 chunks/000002 meta.json index

tombstone

- The initial 2-hour blocks are merged in the background to form longer blocks
- This is called compaction

directory 1

2h of data:

chunks/000001 chunks/000002

directory 2+3

4h of data:

chunks/000001 chunks/000002

Block characteristics:

- A block on the filesystem is a directory with chunks
- You can see each block as a fully independent database containing all time series for the window
- Every block of data, except the current block, is **immutable** (no changes can be made)
- These non-overlapping blocks are actually a horizontal partitioning of the ingested time series data

- This horizontal partitioning gives a lot of benefits:
 - When querying, the blocks not in the time range can be skipped
 - When completing a block, data only needs to be added, and not modified (avoids write-amplification)
 - · Recent data is kept in memory, so can be queried quicker
 - Deleting old data is only a matter of deleting directories on the filesystem

• Compaction:

- When querying, blocks have to be merged together to be able to calculate the results
- Too many blocks could cause **too much merging overhead**, so blocks are compacted
 - 2 blocks are merged and form a newly created (often larger) block
 - Compaction can also modify data: dropping deleted data or restructuring the chunks to increase the query performance

- The index:
 - Having horizontal partitioning already makes most queries quicker, but not those that need to go through all the data to get the result
 - The index is an inverted index to provide better query performance, also in cases where all data needs to be queried
 - Each series is assigned a unique ID (e.g. ID 1, 2, and 3)
 - The index will contain an inverted index for the labels, for example for label env=production, it'll have 1 and 3 as IDs if those series contain the label env=production

- What about Disk size?
- On average, Prometheus needs 1-2 bytes per sample
- You can use the following formula to calculate the disk space needed:

needed_disk_space = retention_time_seconds * ingested_samples_per_second * bytes_per_sample

- How to reduce disk size?
 - You can increase the scrape interval, which will get you less data
 - You can decrease the targets or series you scrape
 - Or you can can reduce the retention (how long you keep the data)

--storage.tsdb.retention: This determines when to remove old data. Defaults to 15d.

References

 To read the full story of Prometheus time series database, read the blog post from Fabian Reinartz at https://fabxc.org/tsdb/

Prometheus

Security

Security

- At the moment Prometheus doesn't offer any support for authentication or encryption (TLS) on the server components
 - They argue that they're focussing on building a monitoring solution, and want to avoid having to implement complex security features
 - You can still enable authentication and TLS, using a reverse proxy
- This is only valid for server components, prometheus can scrape TLS and authentication enabled targets
 - See tls_config in the prometheus configuration to configure a CA certificate, user certificate and user key
 - You'd still need to setup a reverse proxy for the targets itself

Demo

Prometheus TLS and authentication

Demo

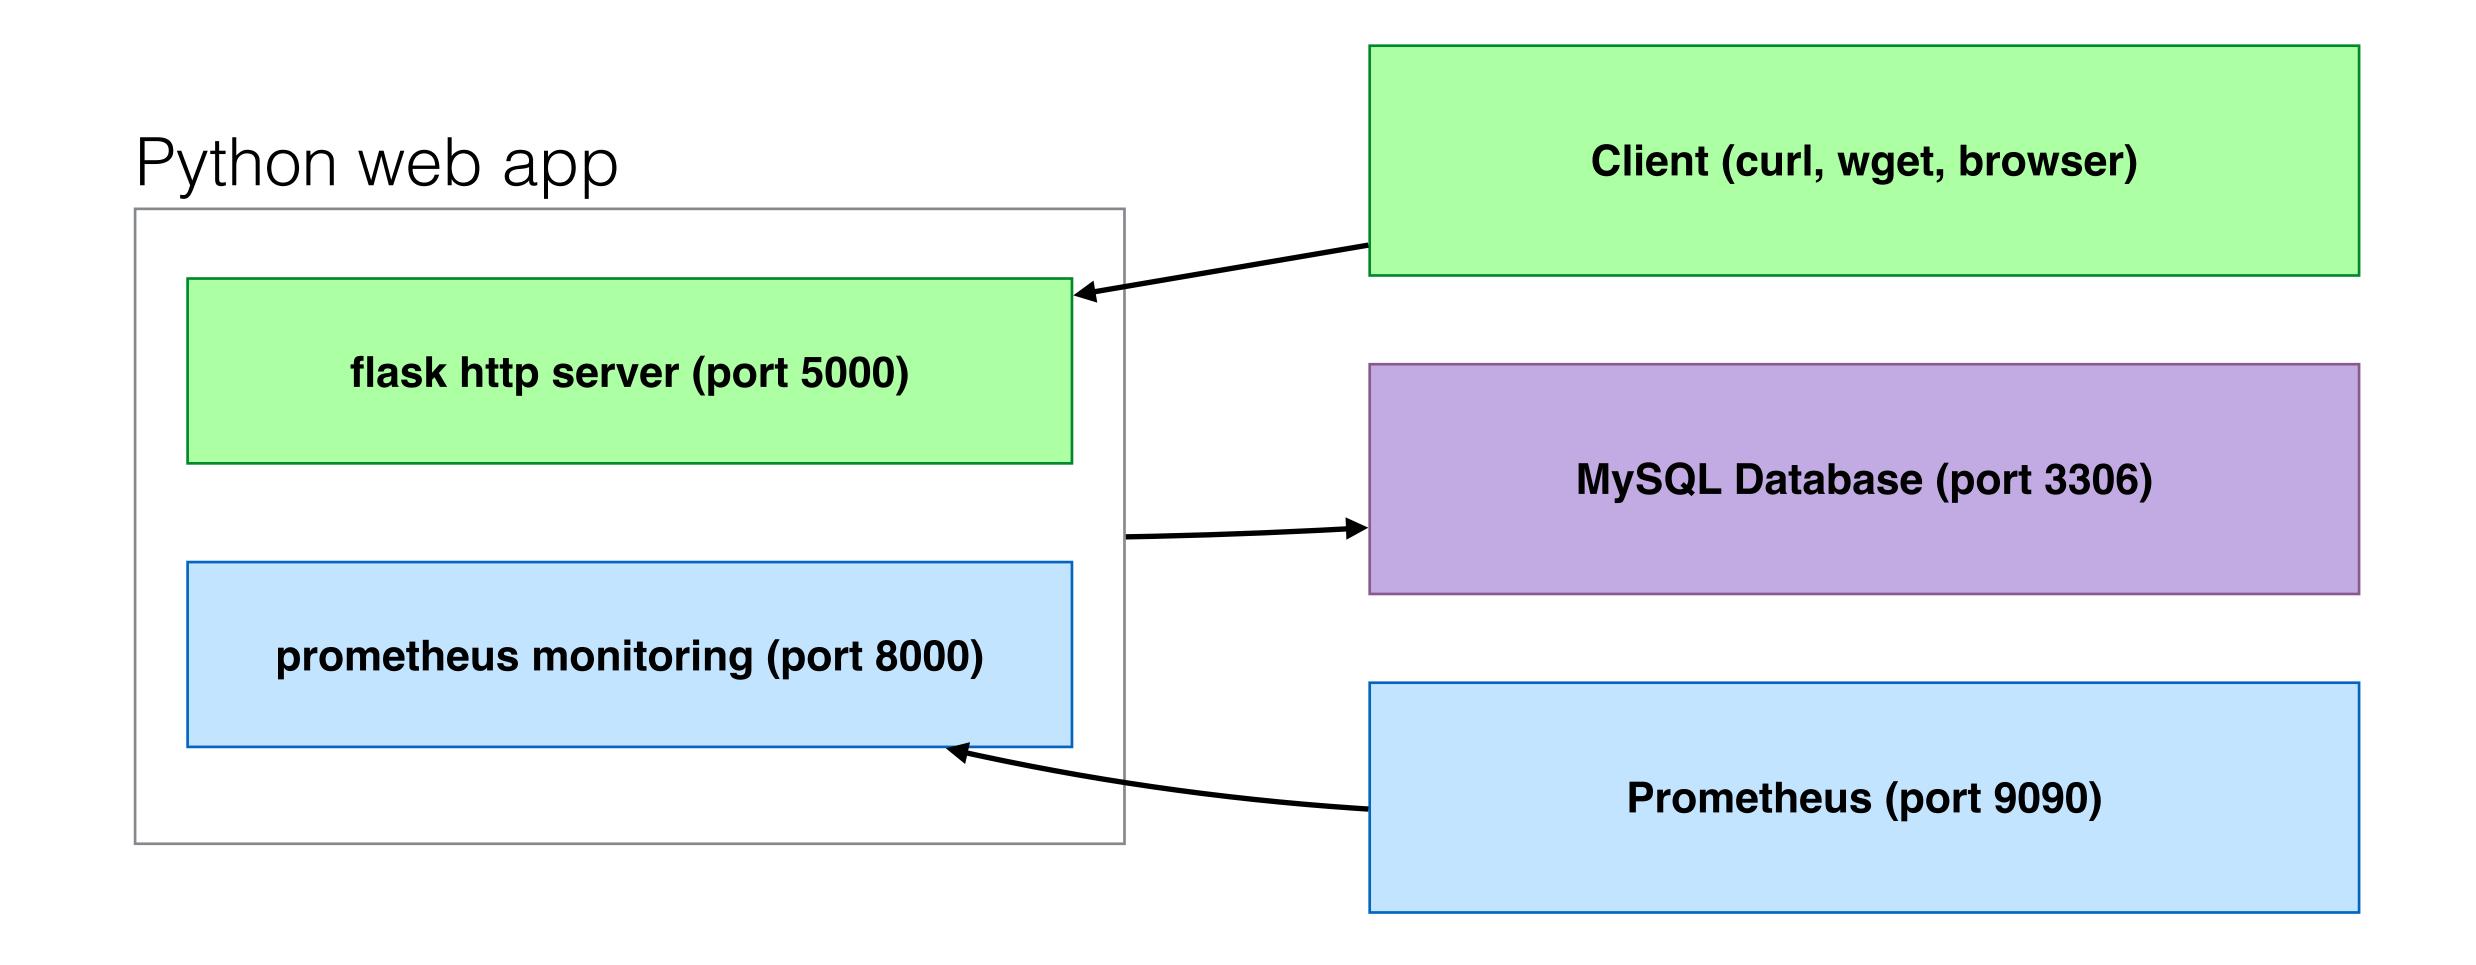
Prometheus mutual TLS for targets

Prometheus Use Cases

Prometheus with python-flask and MySQL

- I'm going to integrate prometheus monitoring with a web application based on python
 - I'll use the official prometheus_client library for Python
 - Flask is the web framework I'm going to use
 - It will create an http server and I'll able to configure routes (e.g. /query)
 - I'll use mysqlclient for python to query a MySQL database
 - I'll include one normal query and one "bad behaving" query that will take between 0 and 10 seconds to execute

The web app



- I'm going to use the Counter and the Histogram metric types to capture the data:
 - A Counter to capture the amount of times an http endpoint is hit + to capture the amount of times a MySQL query is executed
 - The value of the **Counter must always increase**, that's why you should take the Counter type for these types of data
 - A Histogram to capture the latency of the HTTP requests and the MySQL Queries
 - A Histogram samples observations (like latencies) and counts them in configurable buckets. It also provides a sum of all observed values.
 - The default buckets are intended to cover a typical web/rpc request from milliseconds to seconds

This is how I'm going to define the data types in Python for Prometheus:

This is how we can calculate the latency of a query:

```
start_time = time.time()

sql = "select * from table"
# do the query

query_latency = time.time() - start_time

MYSQL_REQUEST_LATENCY.labels(sql[:50]).observe(query_latency)

MYSQL_REQUEST_COUNT.labels(sql[:50]).inc()
```

Demo

Monitor a web application with Prometheus

Demo

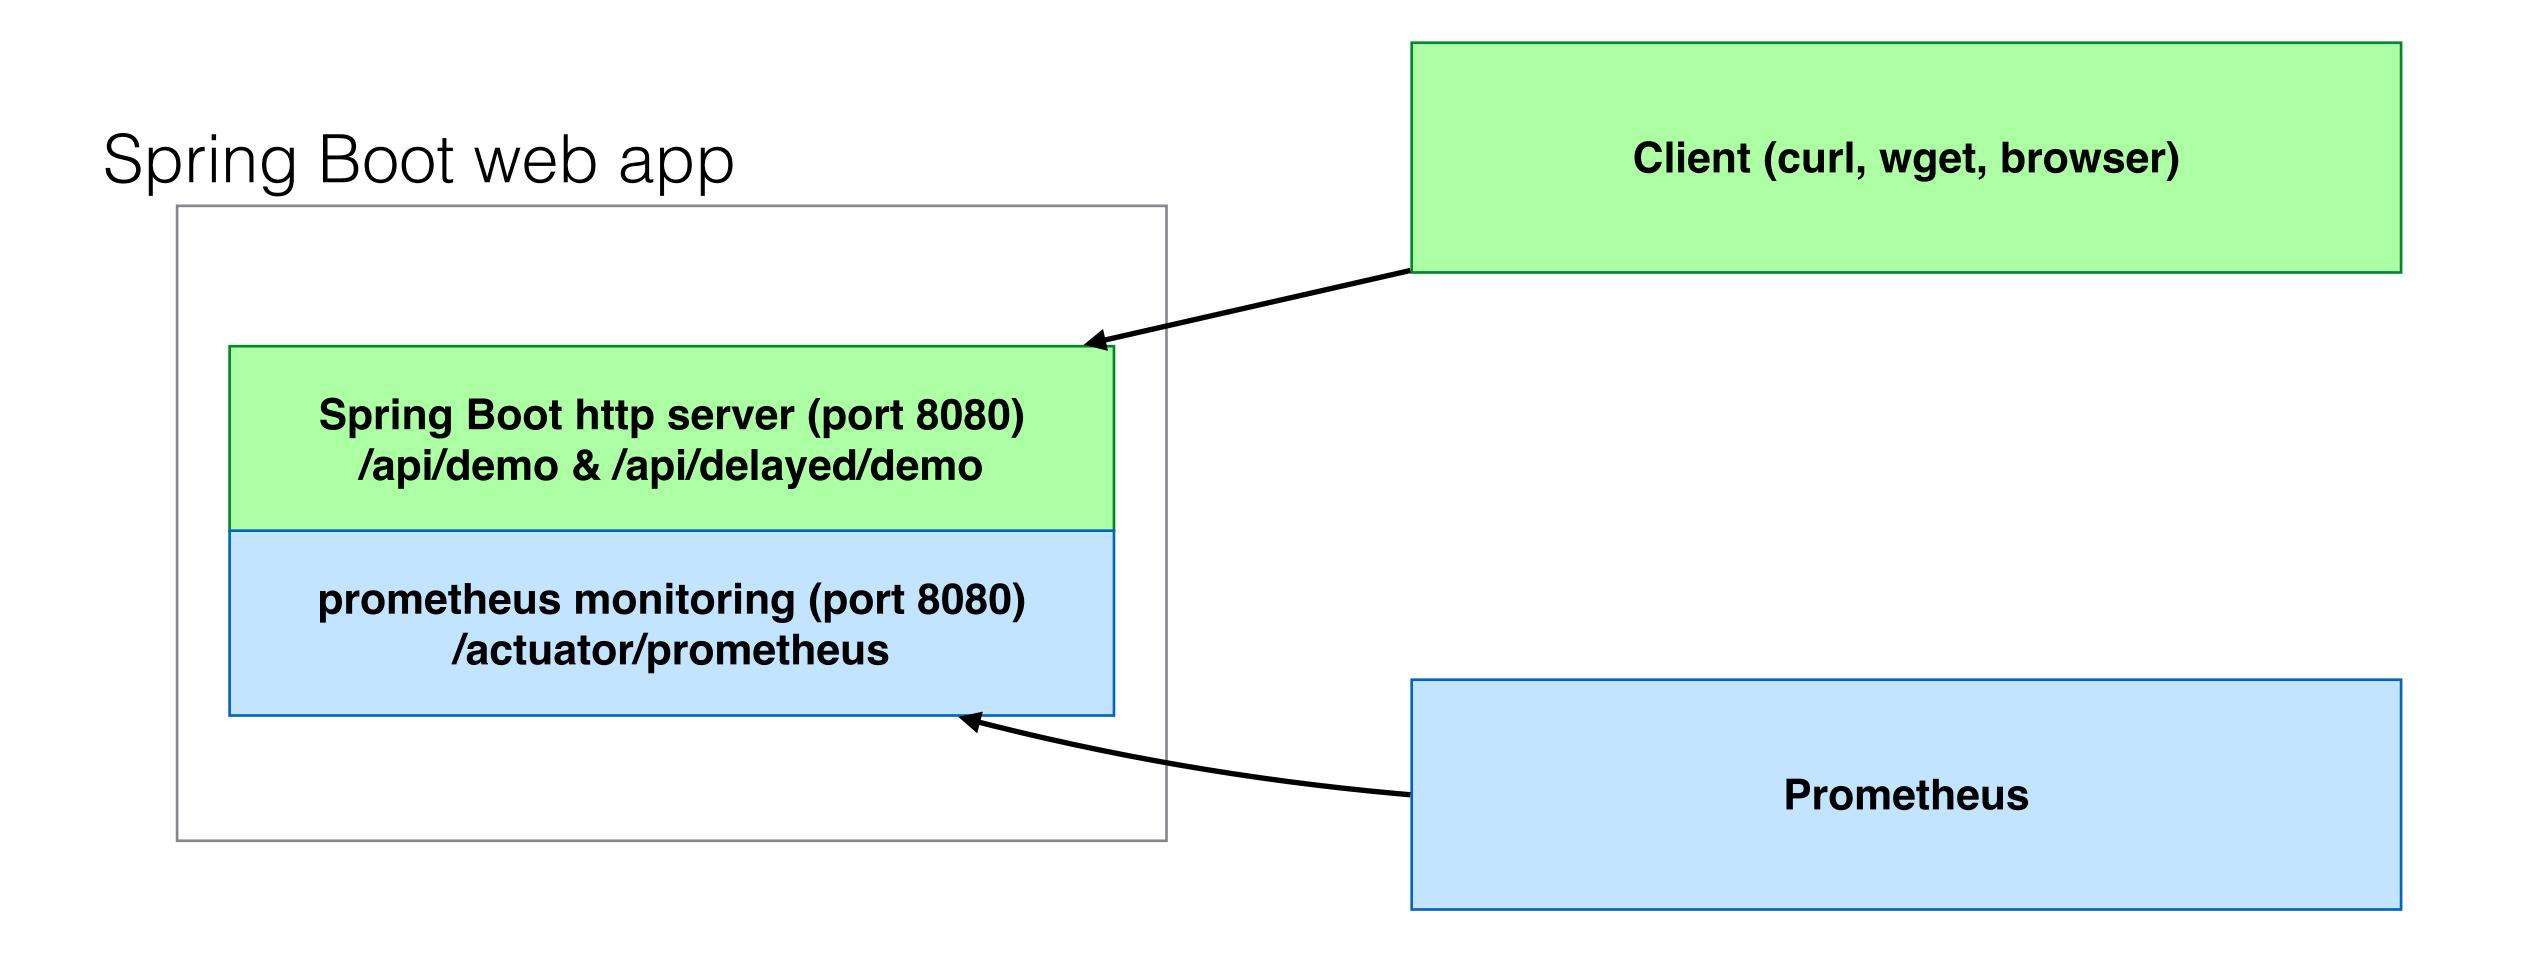
Monitor a web application with Prometheus - apdex score

Prometheus with Java Spring boot

- Introduction
 - Application's health + Metrics
 - Notice unwanted behaviour
 - Monoliths as well as microservices
 - Crucial in microservices architecture
 - "To measure is to know"

- We are going to integrate prometheus monitoring with a web application based on Java Spring Boot
 - We will use following:
 - Spring Boot
 - Spring Boot Actuator
 - Micrometer
 - We are also going to do a demo with an example.

The web app



- Spring Boot Actuator
 - Sub-Project of Spring boot
 - Production ready endpoints
 - /actuator is the common prefix of these endpoints
 - Protected by default
 - Adjustable in application.properties
 - Expose all: management.endpoints.web.exposure.include=*

Micrometer

- Vendor-Neutral application metrics facade
- Support for Prometheus and many others:
 - AWS Cloudwatch, Datadog, InfluxDB/Telegraf, New Relic, ...
- Transforms /actuator/metrics data into data your monitoring system understands
- Only a vendor-specific micrometer dependency in your application is required

Monitoring a web app

Micrometer

pom.xml example

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-core</artifactId>
</dependency>
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-registry-prometheus</artifactId>
</dependency>
```

Monitoring a web app

- Spring Boot
 - Code example

```
import io.micrometer.core.instrument.Metrics;
private Counter runCounter = Metrics.counter("runCounter");
@GetMapping("/api/demo")
@Timed
public String apiUse() throws InterruptedException {
 runCounter.increment();
  log.info("Hello world app accessed on /api/demo");
  return "Hello world";
```

Demo

Monitor and instrument a Spring boot application

Grafana

Grafana Provisioning

Grafana Provisioning

- In one of the first lectures I showed you how to setup Grafana using the UI
- Rather than using the UI, you can also use yaml and json files to provision Grafana with datasources and dashboards
- This is a much more powerful way of using Grafana, as you can test new dashboards first on a dev / test server, then import the newly created dashboards to production
 - You can do the import manually through the UI, or using yaml and json files
 - When using files, you can keep files within version control to keep changes, revisions and backups

Grafana Provisioning

The configuration of Grafana is all kept in /etc/grafana:

/etc/grafana/:

```
-rw-r---- 1 root grafana 14K Jul 17 12:30 grafana.ini
-rw-r---- 1 root grafana 3.4K Jul 17 12:30 ldap.toml
drwxr-xr-x 4 root grafana 4.0K Jul 17 13:15 provisioning/
/etc/grafana/provisioning/:
```

drwxr-xr-x 2 root grafana 4.0K Jul 17 14:56 dashboards/drwxr-xr-x 2 root grafana 4.0K Jul 17 15:34 datasources/

• The data is kept in /var/lib/grafana:

/var/lib/grafana:

```
drwxr-xr-x 2 root root 4.0K Jul 17 15:47 dashboards/
-rw-r---- 1 grafana grafana 500K Jul 17 15:48 grafana.db
drwxr-x--- 2 grafana grafana 4.0K Jul 17 12:31 plugins/
drwx----- 5 grafana grafana 4.0K Jul 17 12:40 sessions/
```

Grafana Provisioning

You can change the database & paths in /etc/grafana/grafana.ini

```
[paths]
# Path to where grafana can store temp files, sessions, and the sqlite3 db (if that is used)
;data = /var/lib/grafana
# Directory where grafana can store logs
;logs = /var/log/grafana
# Directory where grafana will automatically scan and look for plugins
;plugins = /var/lib/grafana/plugins
# folder that contains provisioning config files that grafana will apply on startup and while running.
;provisioning = conf/provisioning
[database]
# Either "mysql", "postgres" or "sqlite3", it's your choice
;type = sqlite3
;host = 127.0.0.1:3306
;name = grafana
;user = root
# If the password contains # or ; you have to wrap it with triple quotes. Ex """#password;"""
;password =
```

Demo

Grafana Provisioning

Prometheus Use Cases

Cloudwatch exporter

Use Cases - Cloudwatch

- Cloudwatch exporter
 - Installation
- Configuration (exporter + AWS)
- Charges + measuring them
- Querying metrics

Prometheus Use Cases

- Consul is a distributed, highly available solution providing:
 - A Service Mesh
 - Service Discovery
 - Health checks for your services
 - A Key-Value store
 - Secure Service Communications
 - Multi-datacenter support
- Consul is often deployed in conjunction with Docker

- There are 2 integrations that are interesting to use:
 - 1) Prometheus can **scrape** Consul's metrics and provide you with all sorts of information about your running services
 - Consul provides Service Discovery, so it knows where services are running and what the current state of it is
 - 2) Consul can be integrated within Prometheus to automatically add the services as targets
 - Consul will discover your services, and these can then be automatically added to Prometheus as a target

- In the next demo I'll focus on the **Prometheus integration** with Consul, not really on implementing consul itself
 - I'll show you the **installation of consul**, but not how to integrate consul with your infrastructure (it's out of scope for this Prometheus course)
 - I'll manually register a service to consul, rather than setting up service discovery
 - In production environments, where you have a lot of services, service discovery using consul will allow you to register your services automatically
 - If you are interested in Consul, have a look at the documentation at https://www.consul.io/ and you can find the service registrator for docker at https://github.com/gliderlabs/registrator

Demo

Prometheus Use Cases

EC2 auto discovery

EC2 auto discovery

- Service discovery is the automatic detection of devices and services offered by these devices on a computer network.
- In this Use Case we will:
 - Create prerequisites in AWS (IAM role, Security Groups, EC2 instances)
 - Alter Prometheus config (/etc/prometheus/prometheus.yml)
 - Query the data in Grafana
 - https://github.com/in4it/prometheus-course/blob/master/use-cases/ec2auto-discovery/lab.txt

Prometheus on Kubernetes

Getting Kubernetes metrics